

# Deep learning

## Episode 1

# Basic neural networks



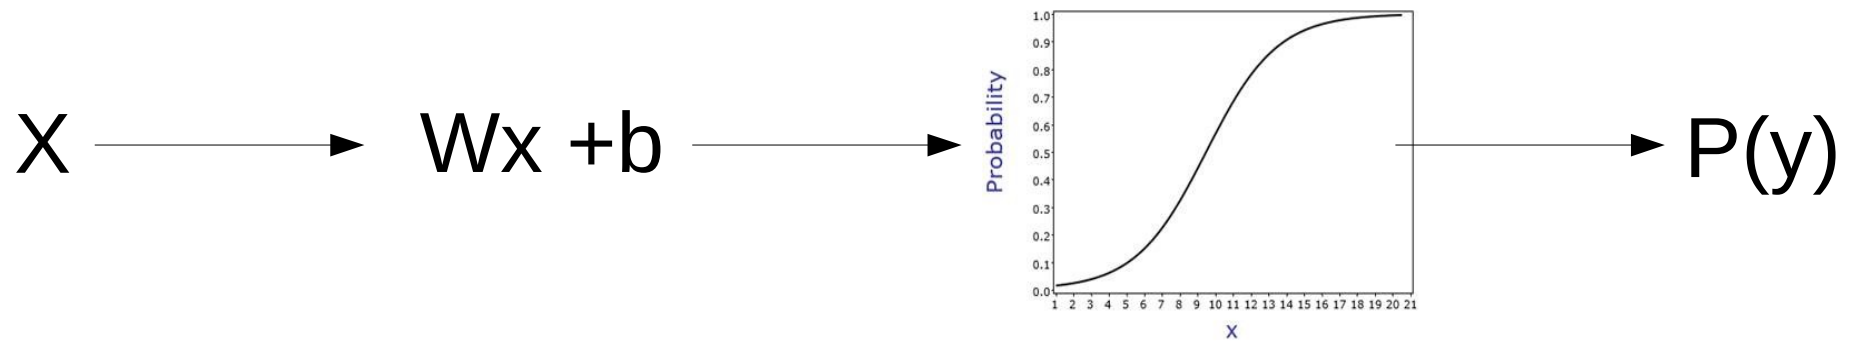
Yandex  
Data Factory

LAMBDA 



**British Hedgehog  
Preservation Society**

# Recap: logistic regression



# Gradient descent

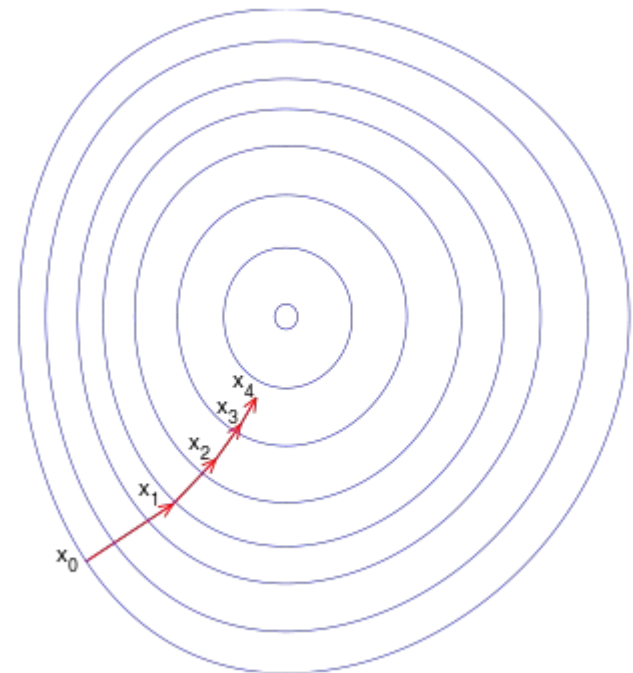
$$y_{pred}(\bar{x}) = \sigma(\bar{w} \cdot \bar{x} + b)$$

$$L(y, y_{pred}) = -(y \cdot \log y_{pred} + (1 - y) \cdot \log(1 - y_{pred}))$$

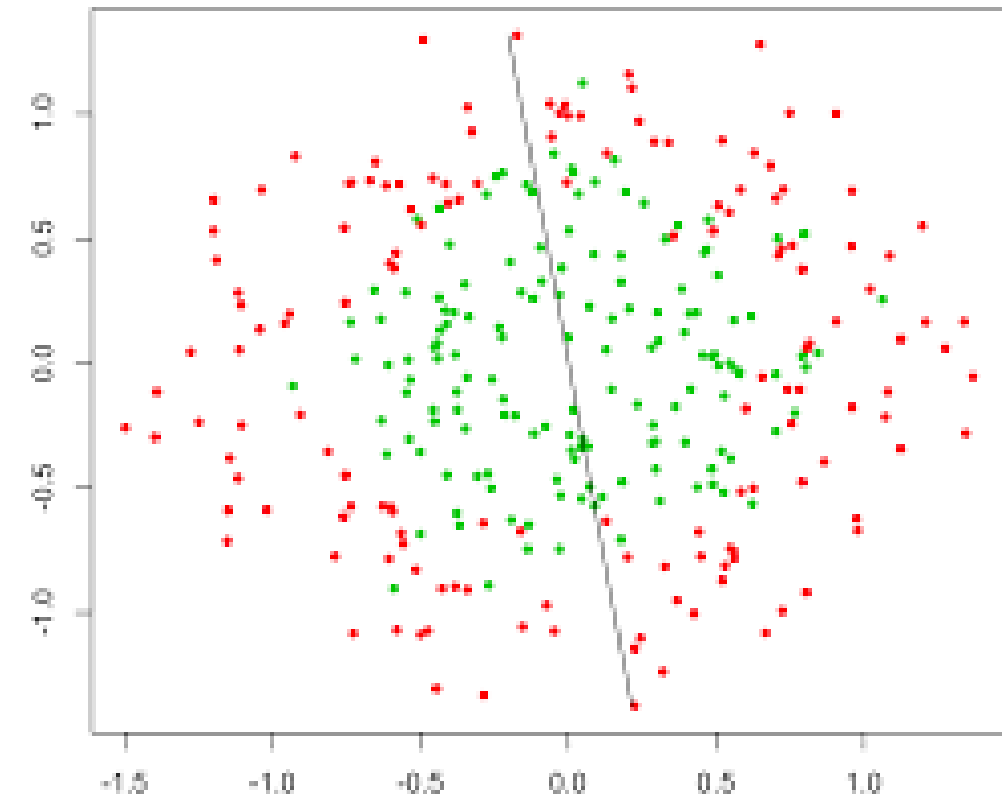
Repeat until convergence

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial L(y, y_{pred})}{\partial \theta_j}$$

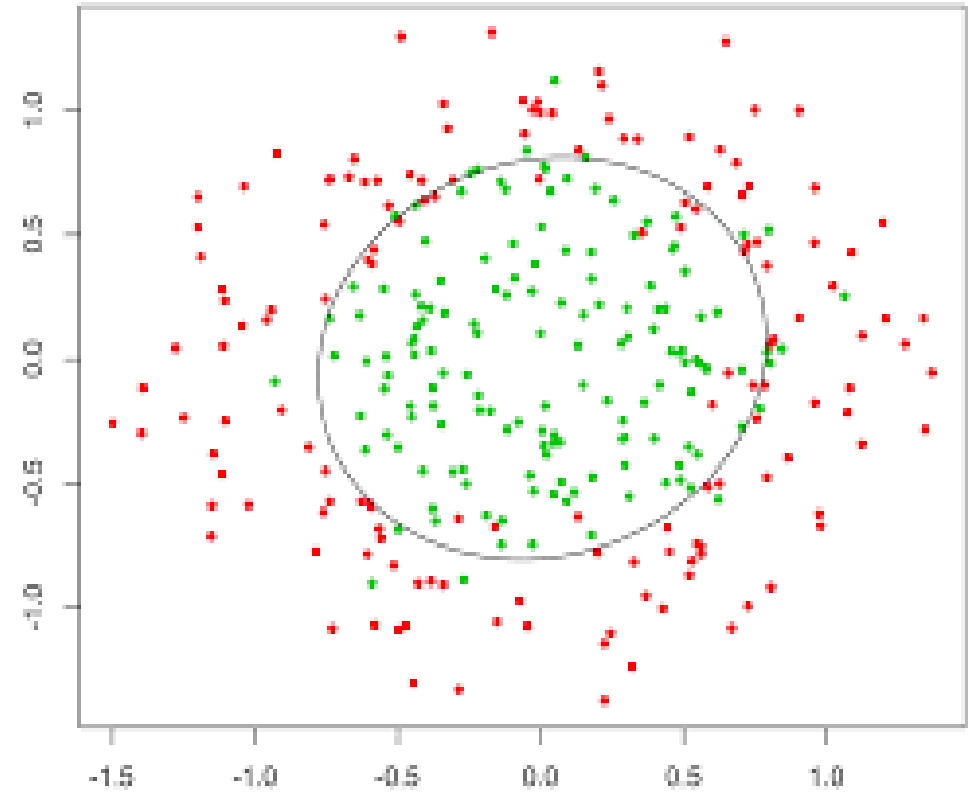
$$\Theta \sim \{w, b\}$$



# Nonlinear dependencies



What we have



What we want

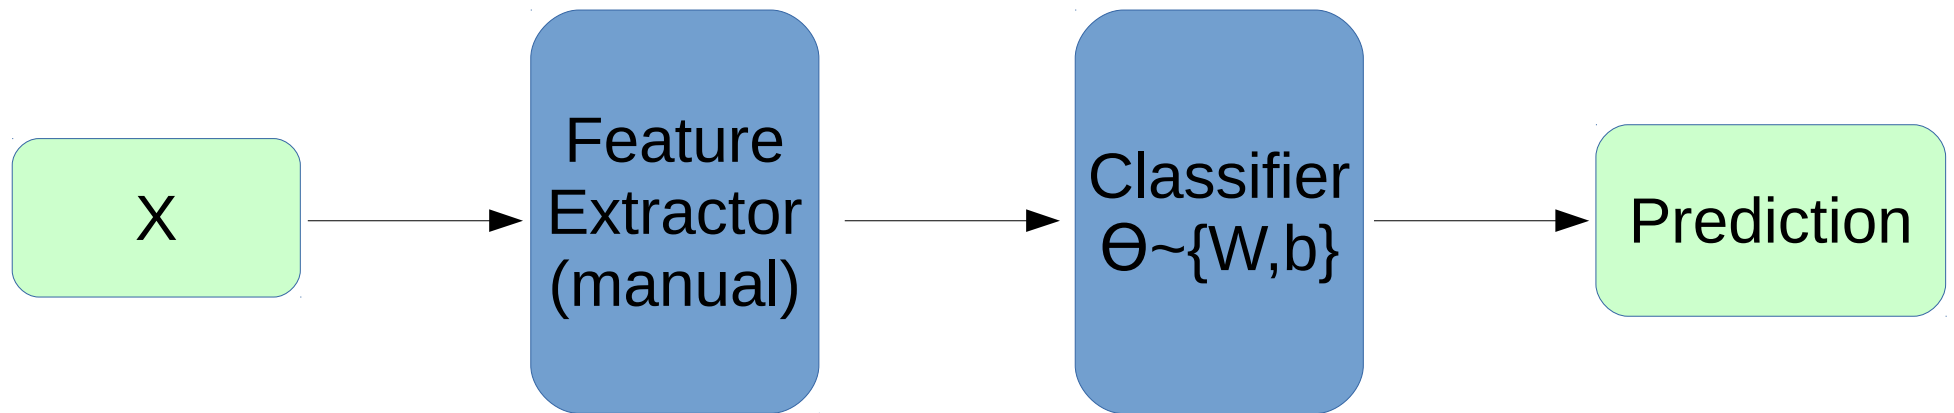
- How to get that?

# Feature extraction

Loss, for example:

$$L(y, y_{pred}) = -(y \cdot \log y_{pred} + (1 - y) \cdot \log(1 - y_{pred}))$$

Model:



Training:

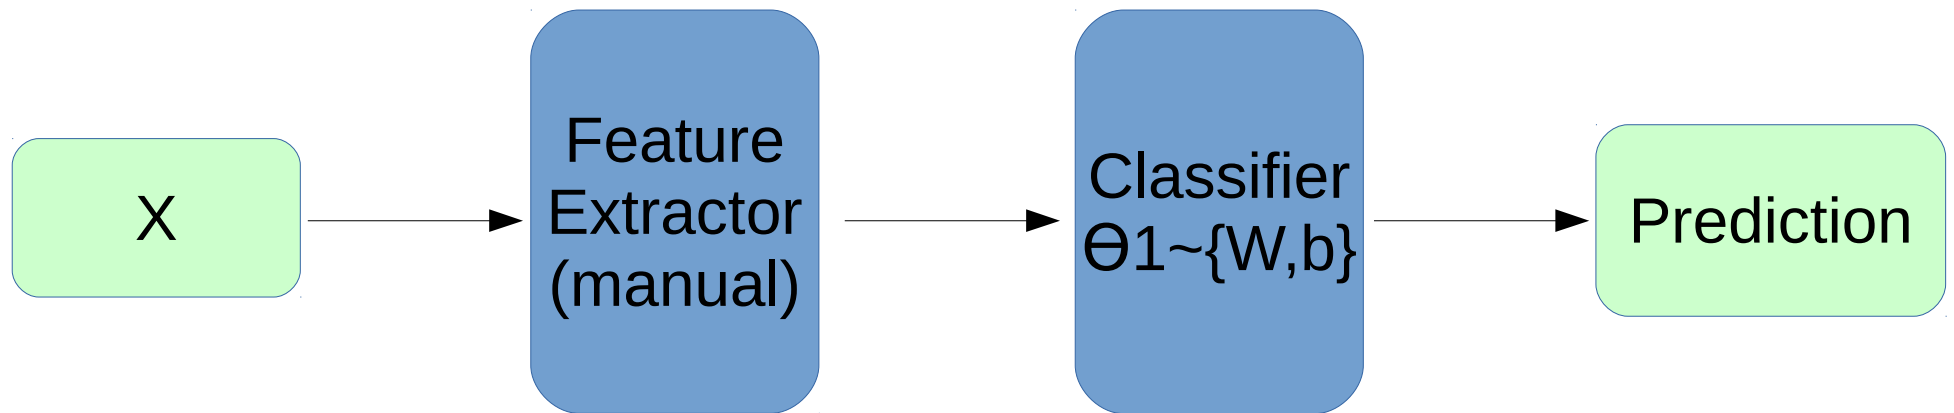
$$\operatorname{argmin}_{\theta} L(y, y_{pred}(X, \theta))$$

# Feature extraction

Loss, for example:

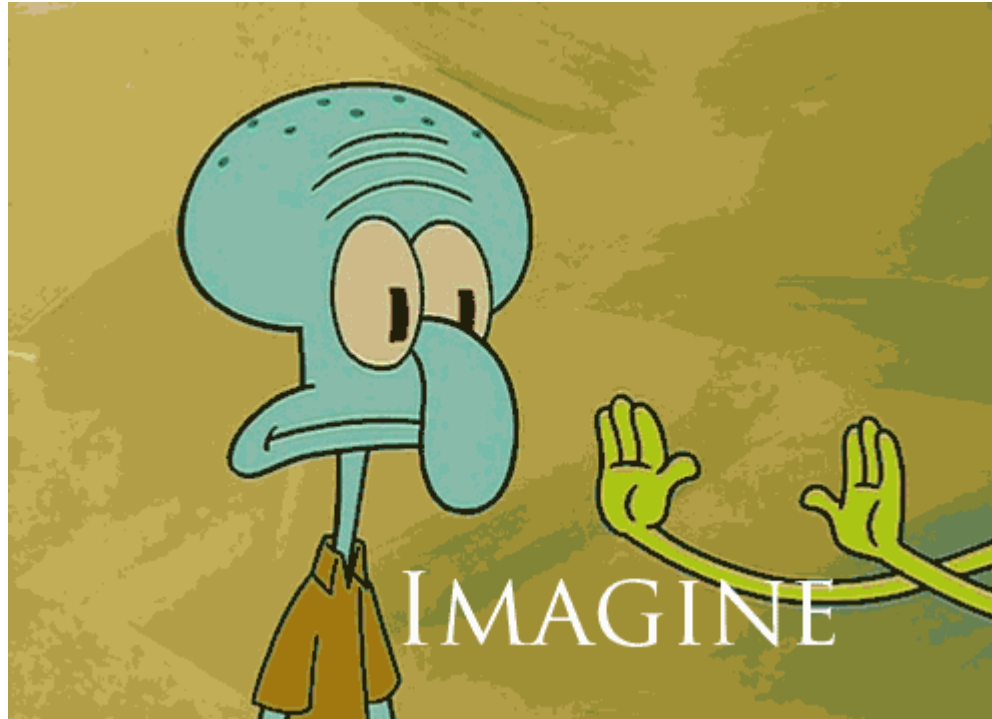
$$L(y, y_{pred}) = -(y \cdot \log y_{pred} + (1 - y) \cdot \log(1 - y_{pred}))$$

Model:



Gradient:

$$\frac{\delta L(y, y_{pred}(X, \theta 1))}{\delta \theta 1}$$



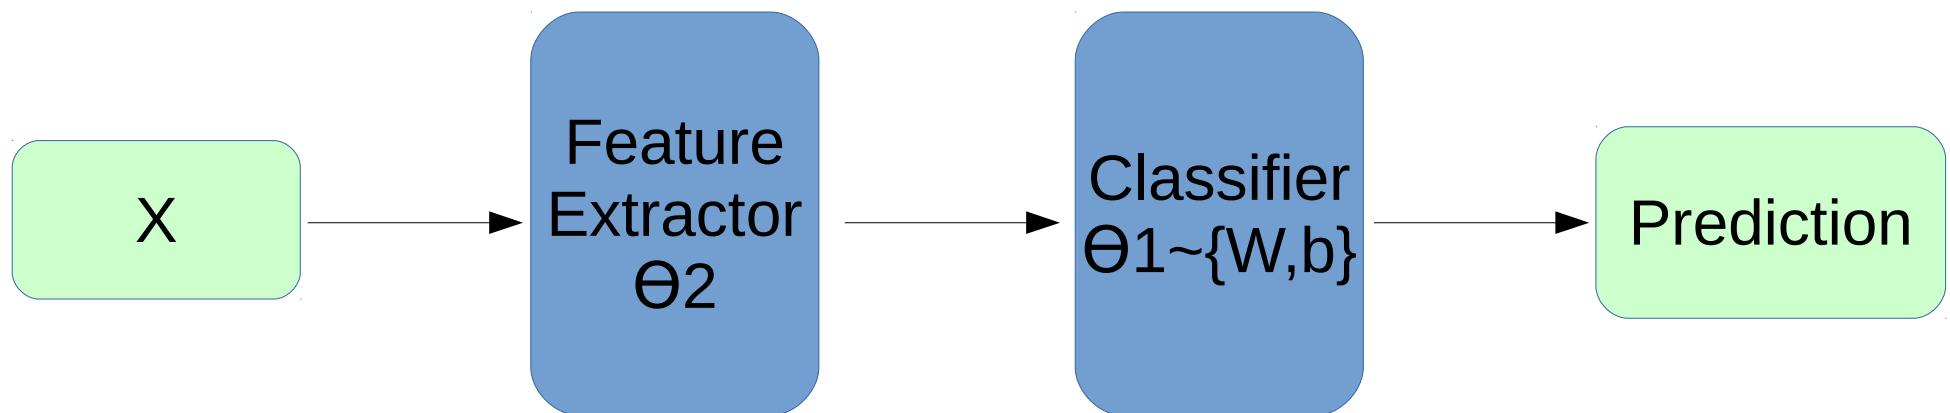
Features would tune to your problem automatically!

# What do we want, exactly?

Loss, for example:

$$L(y, y_{pred}) = -(y \cdot \log y_{pred} + (1 - y) \cdot \log(1 - y_{pred}))$$

Model:



Training:

?

$$\operatorname{argmin}_{\theta_1} L(y, y_{pred}(X, \theta_1, \theta_2))$$

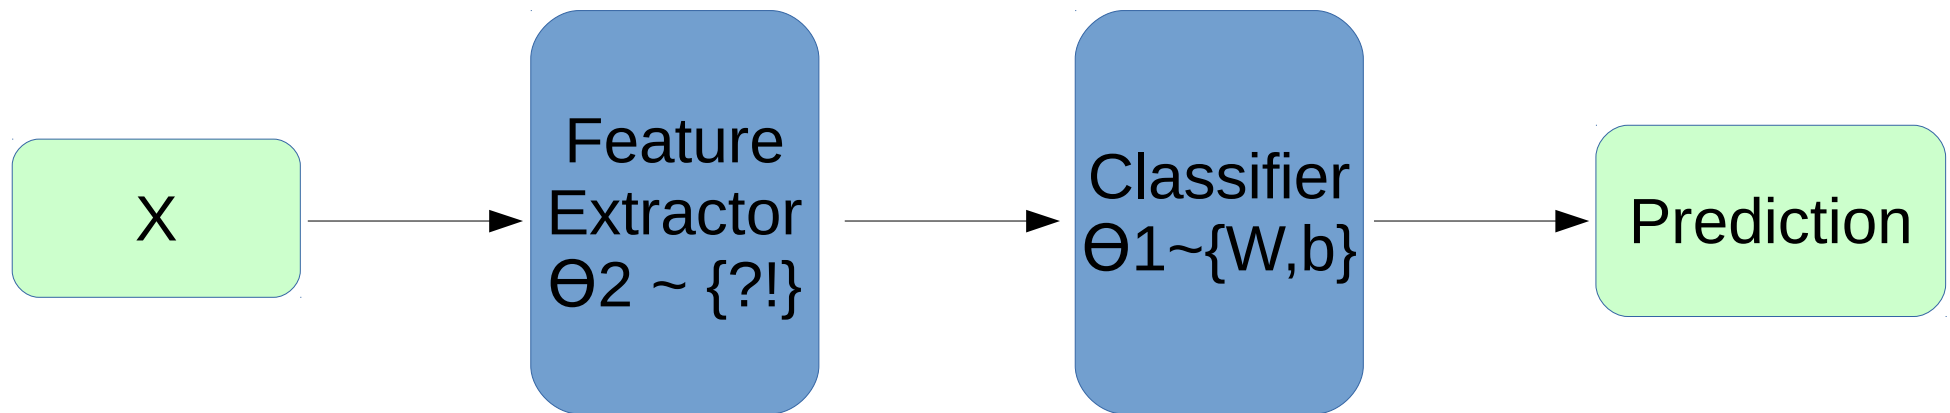


# What do we want, exactly?

Loss, for example:

$$L(y, y_{pred}) = -(y \cdot \log y_{pred} + (1 - y) \cdot \log(1 - y_{pred}))$$

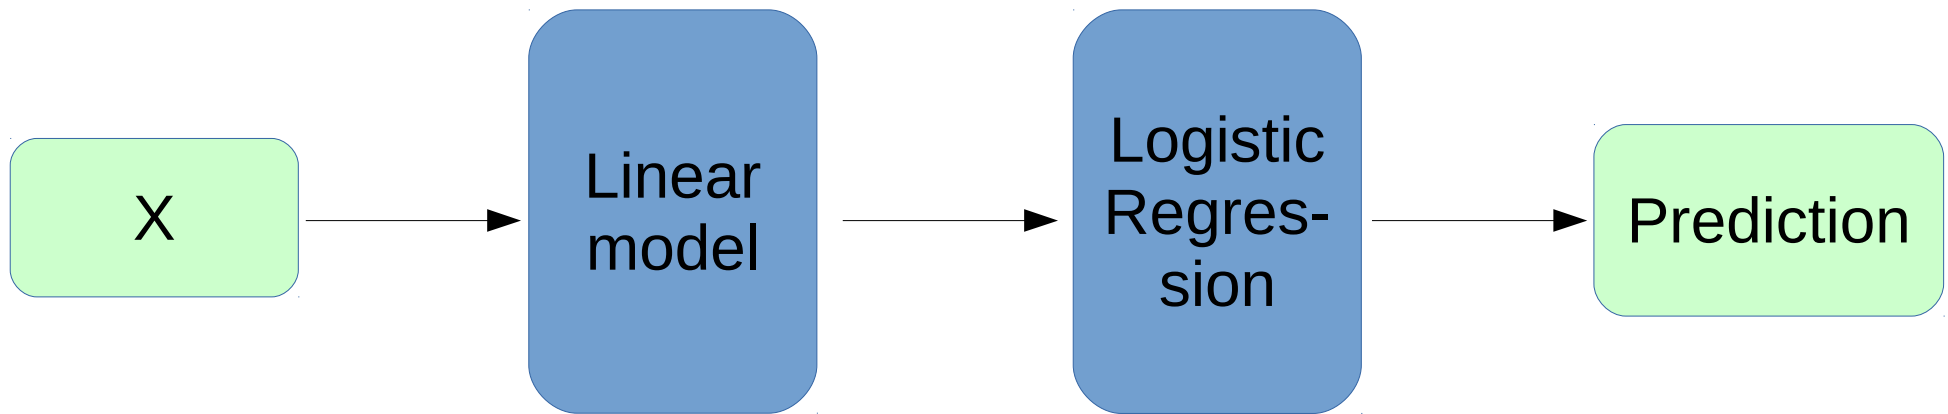
Model:



Gradients:  $\frac{\delta L(y, y_{pred}(X, \theta_1, \theta_2))}{\delta \theta_2}$        $\frac{\delta L(y, y_{pred}(X, \theta_1, \theta_2))}{\delta \theta_1}$

# Try linear

Model:

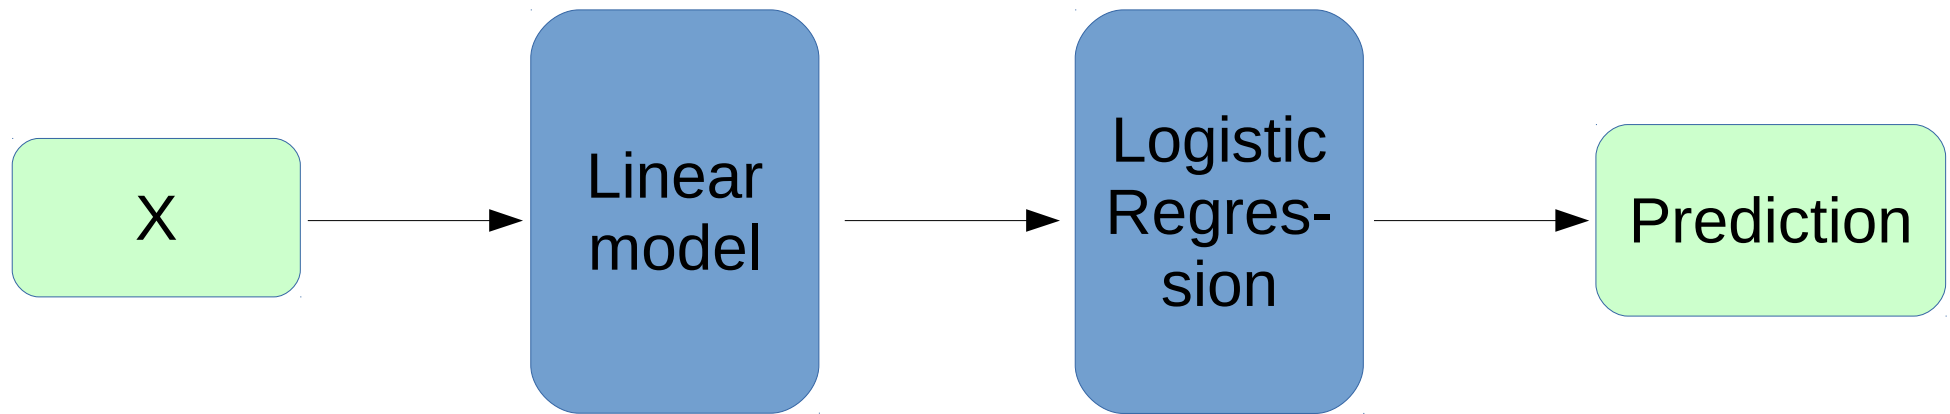


$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

# Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$y_{pred} = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

# Try linear

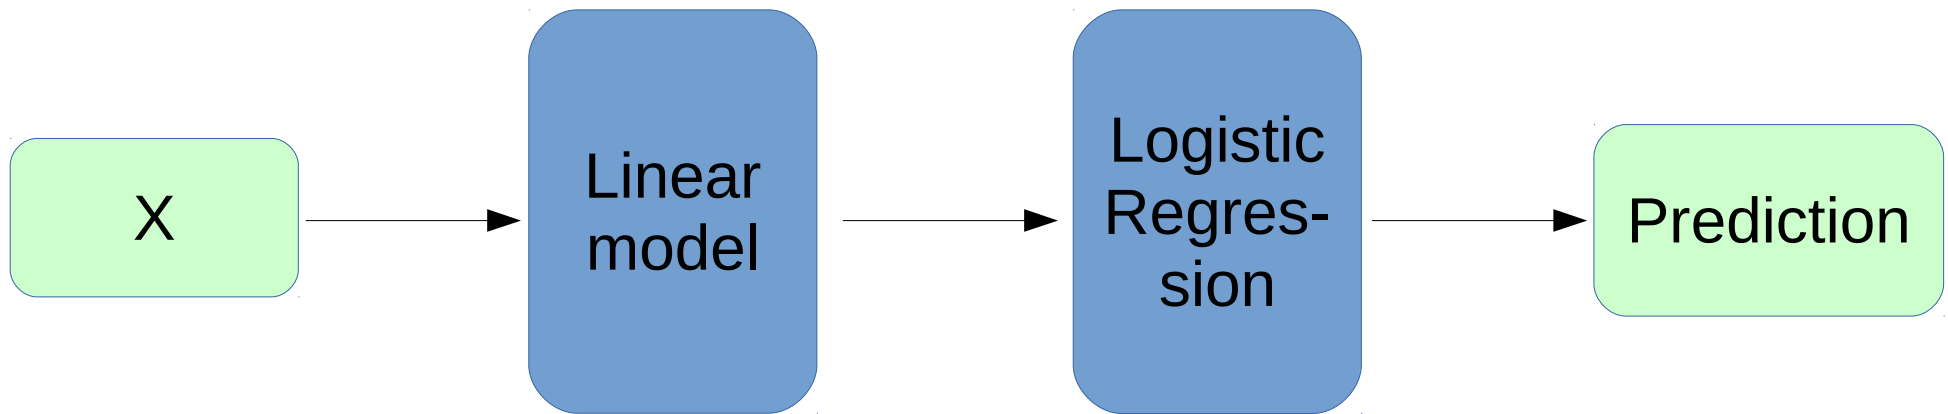
$$y_{pred} = \sigma \left( \sum_j w_j^o \left( \sum_i w_{ij}^h x_i + b_j^h \right) + b^o \right)$$

$$w'_i = \sum_j w_j^o w_{ij}^h \qquad b' = \sum_j w_j^o b_j^h + b^o$$

$$y_{pred} = \sigma \left( \sum_i w'_i x_i + b' \right)$$

# Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

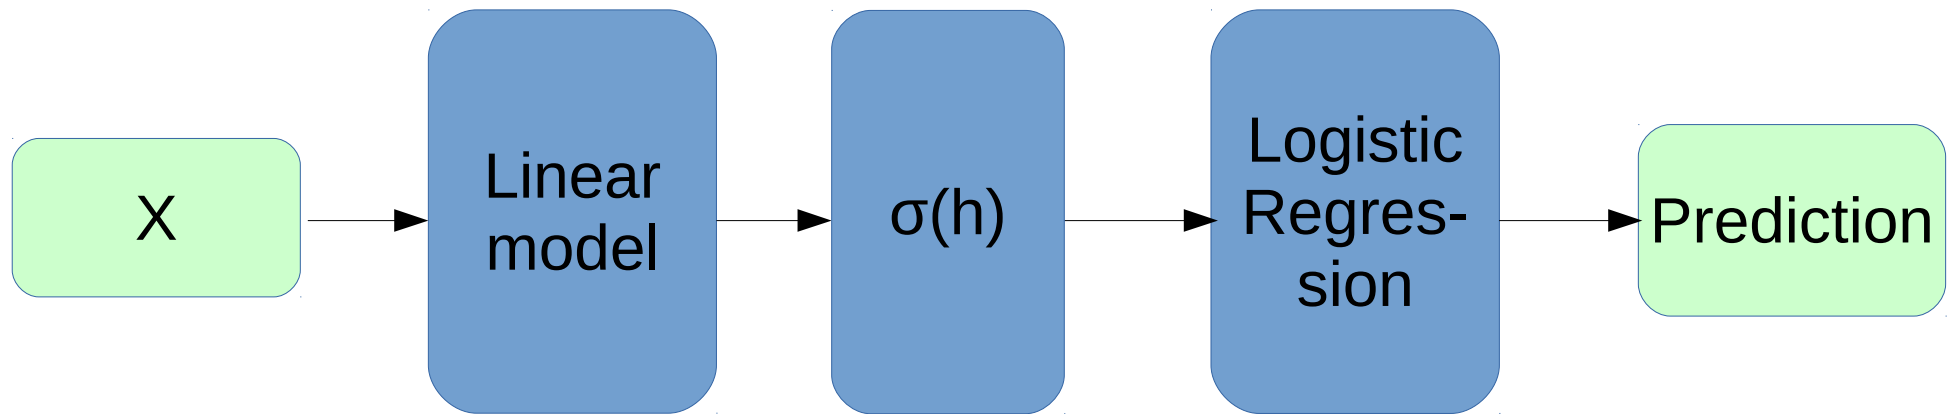
Output:

$$y_{pred} = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

# Nonlinearity

Model:

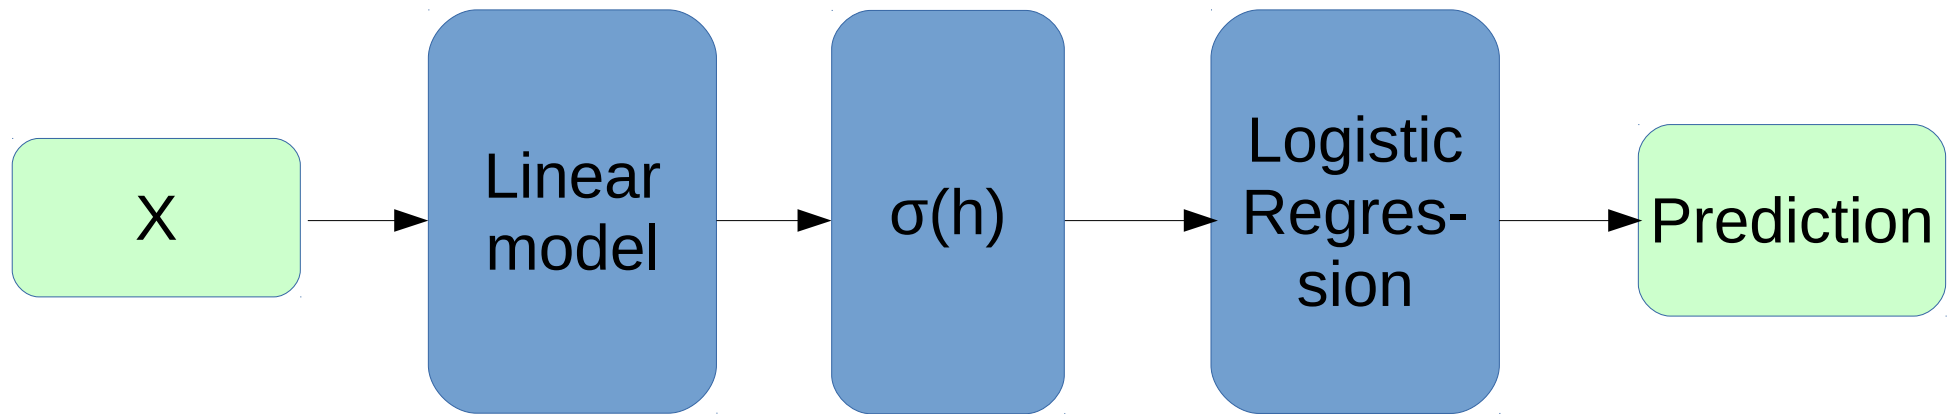


$$h_j = \sigma\left(\sum_{j \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

# Nonlinearity

Model:



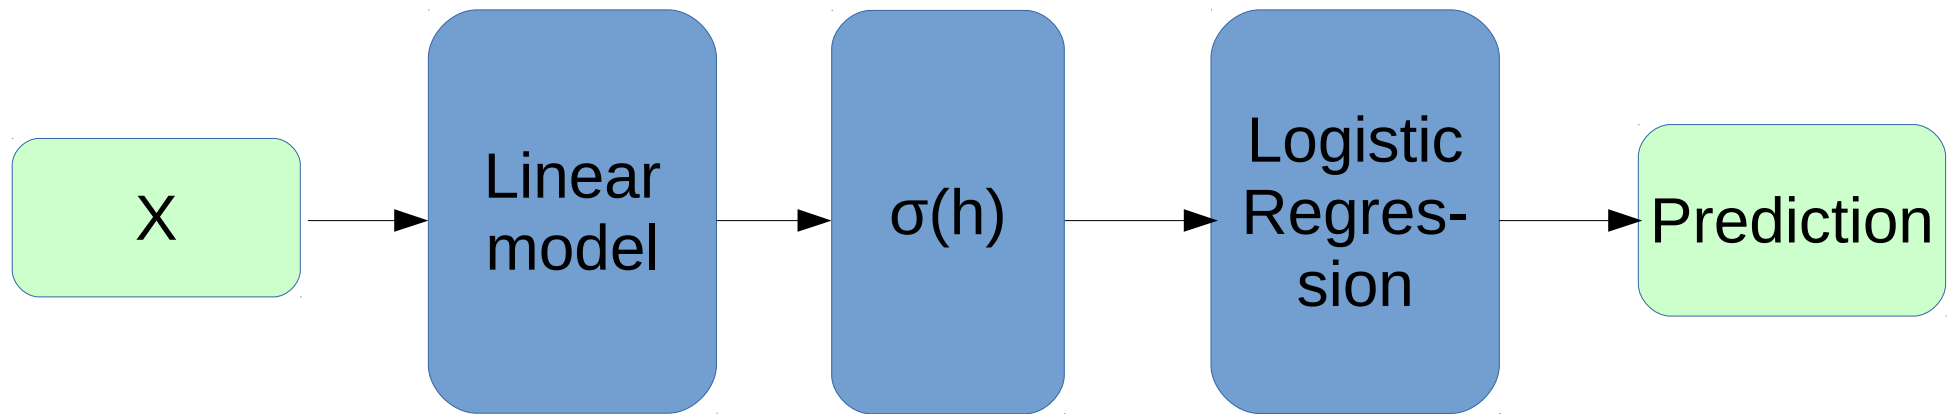
$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right) \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$y_{pred} = \sigma\left(\sum_j w_j^o \sigma\left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

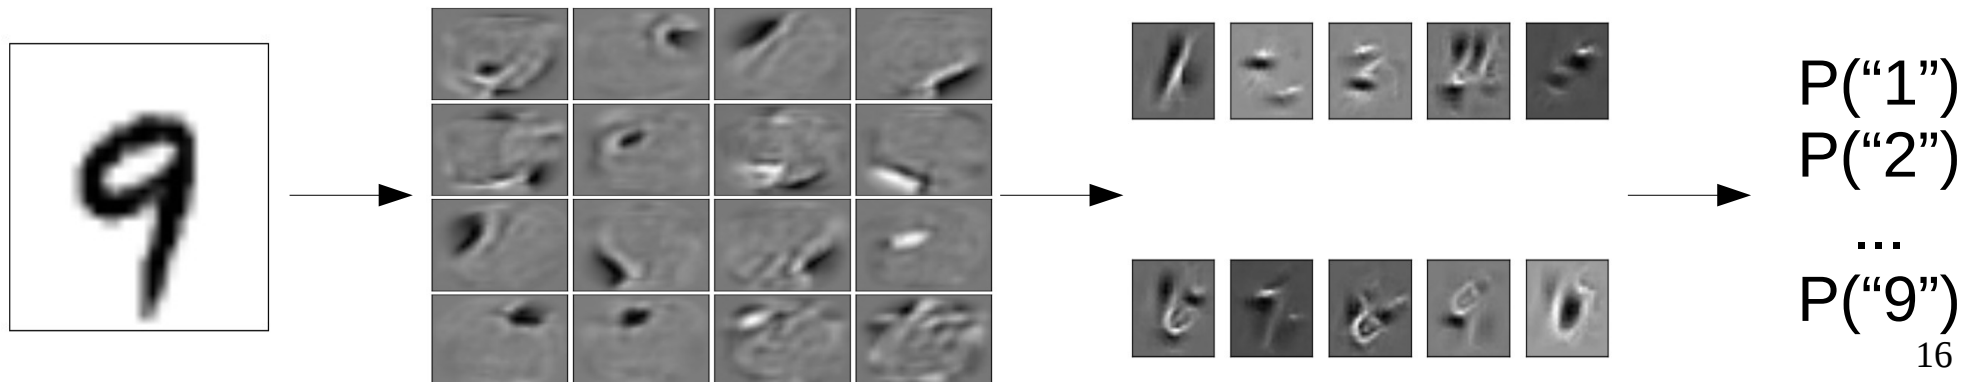
# Nonlinearity

Model:



$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$

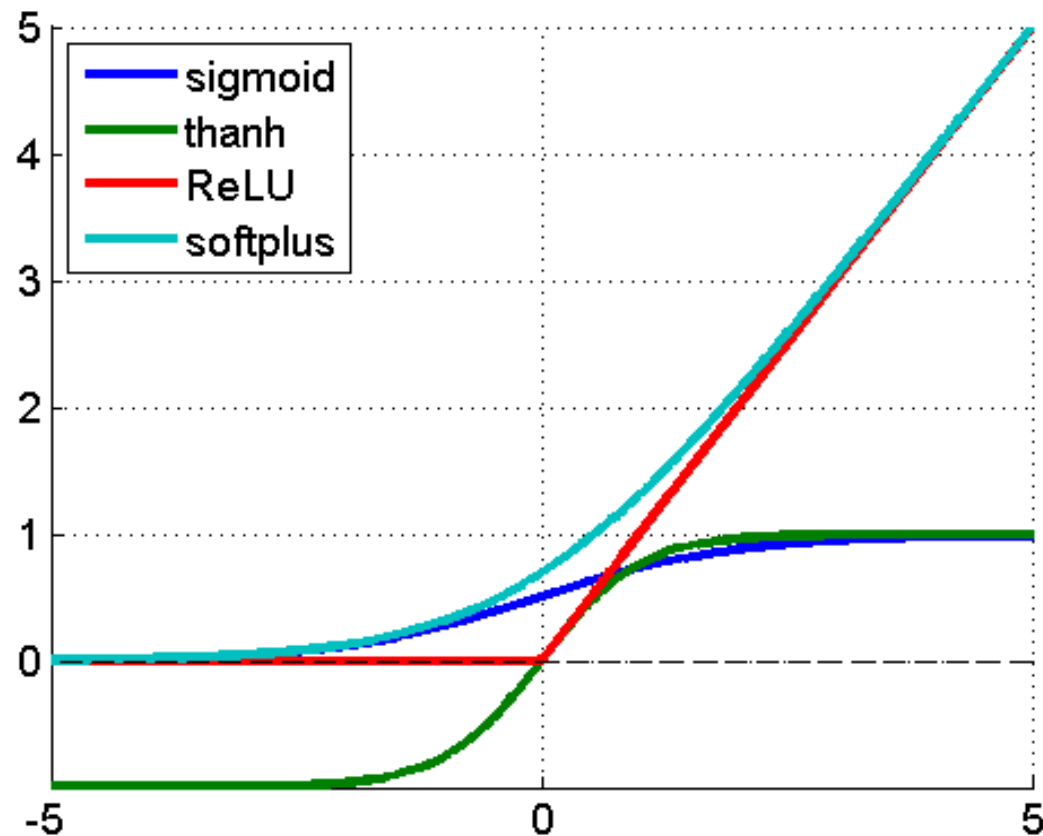
$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$



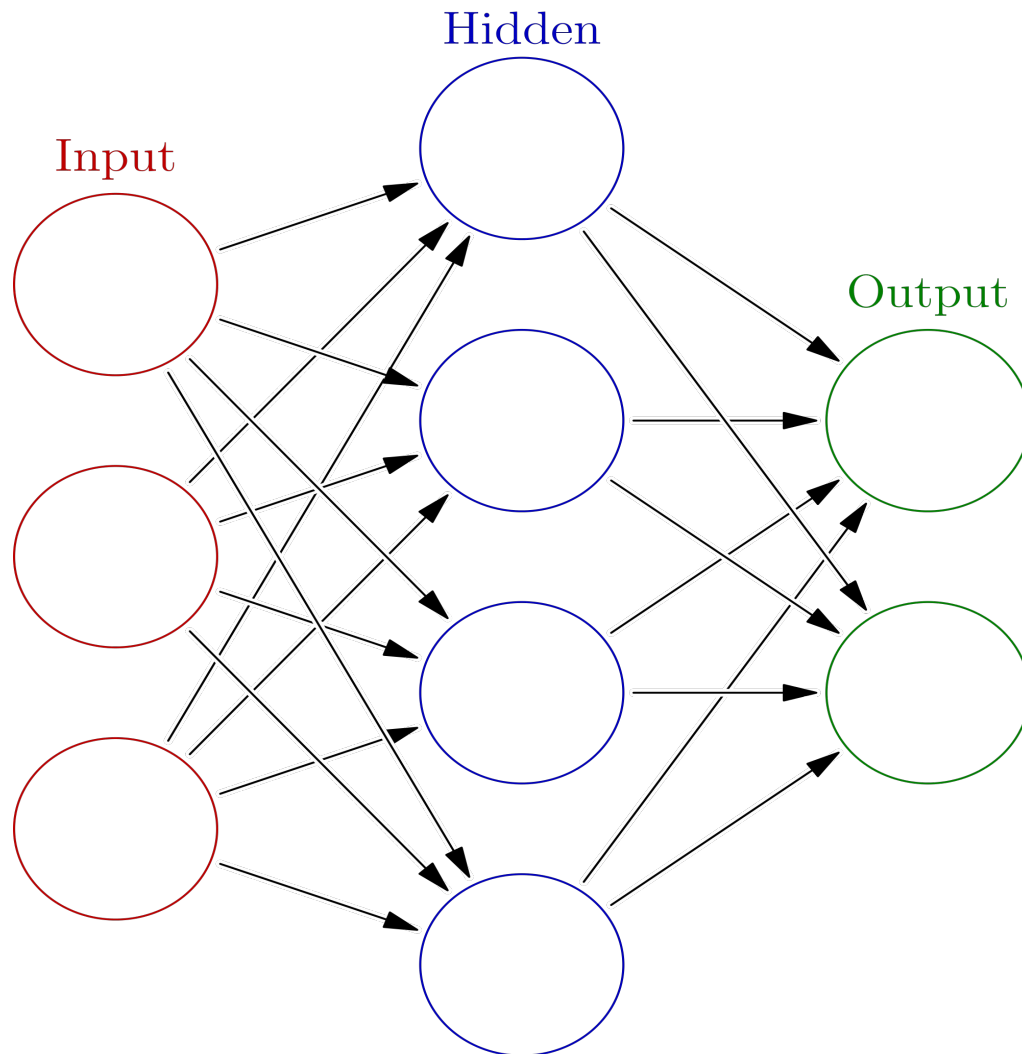


# Nonlinearity

- $f(a) = 1/(1+e^a)$
- $f(a) = \tanh(a)$
- $f(a) = \max(0, a)$
- $f(a) = \log(1+e^x)$

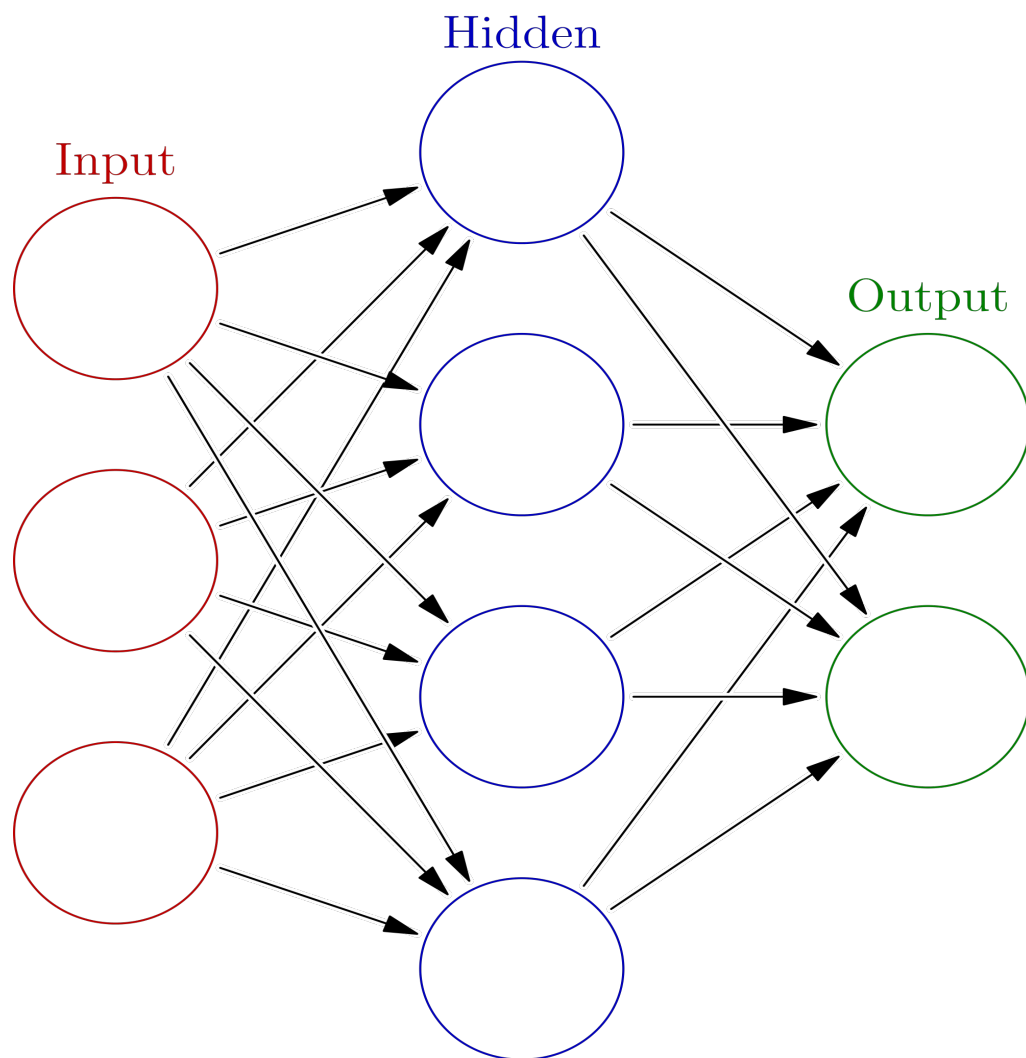


# Initialization, symmetry problem



- Initialize with zeros  
 $W \leftarrow 0$
- What will the first step look like?

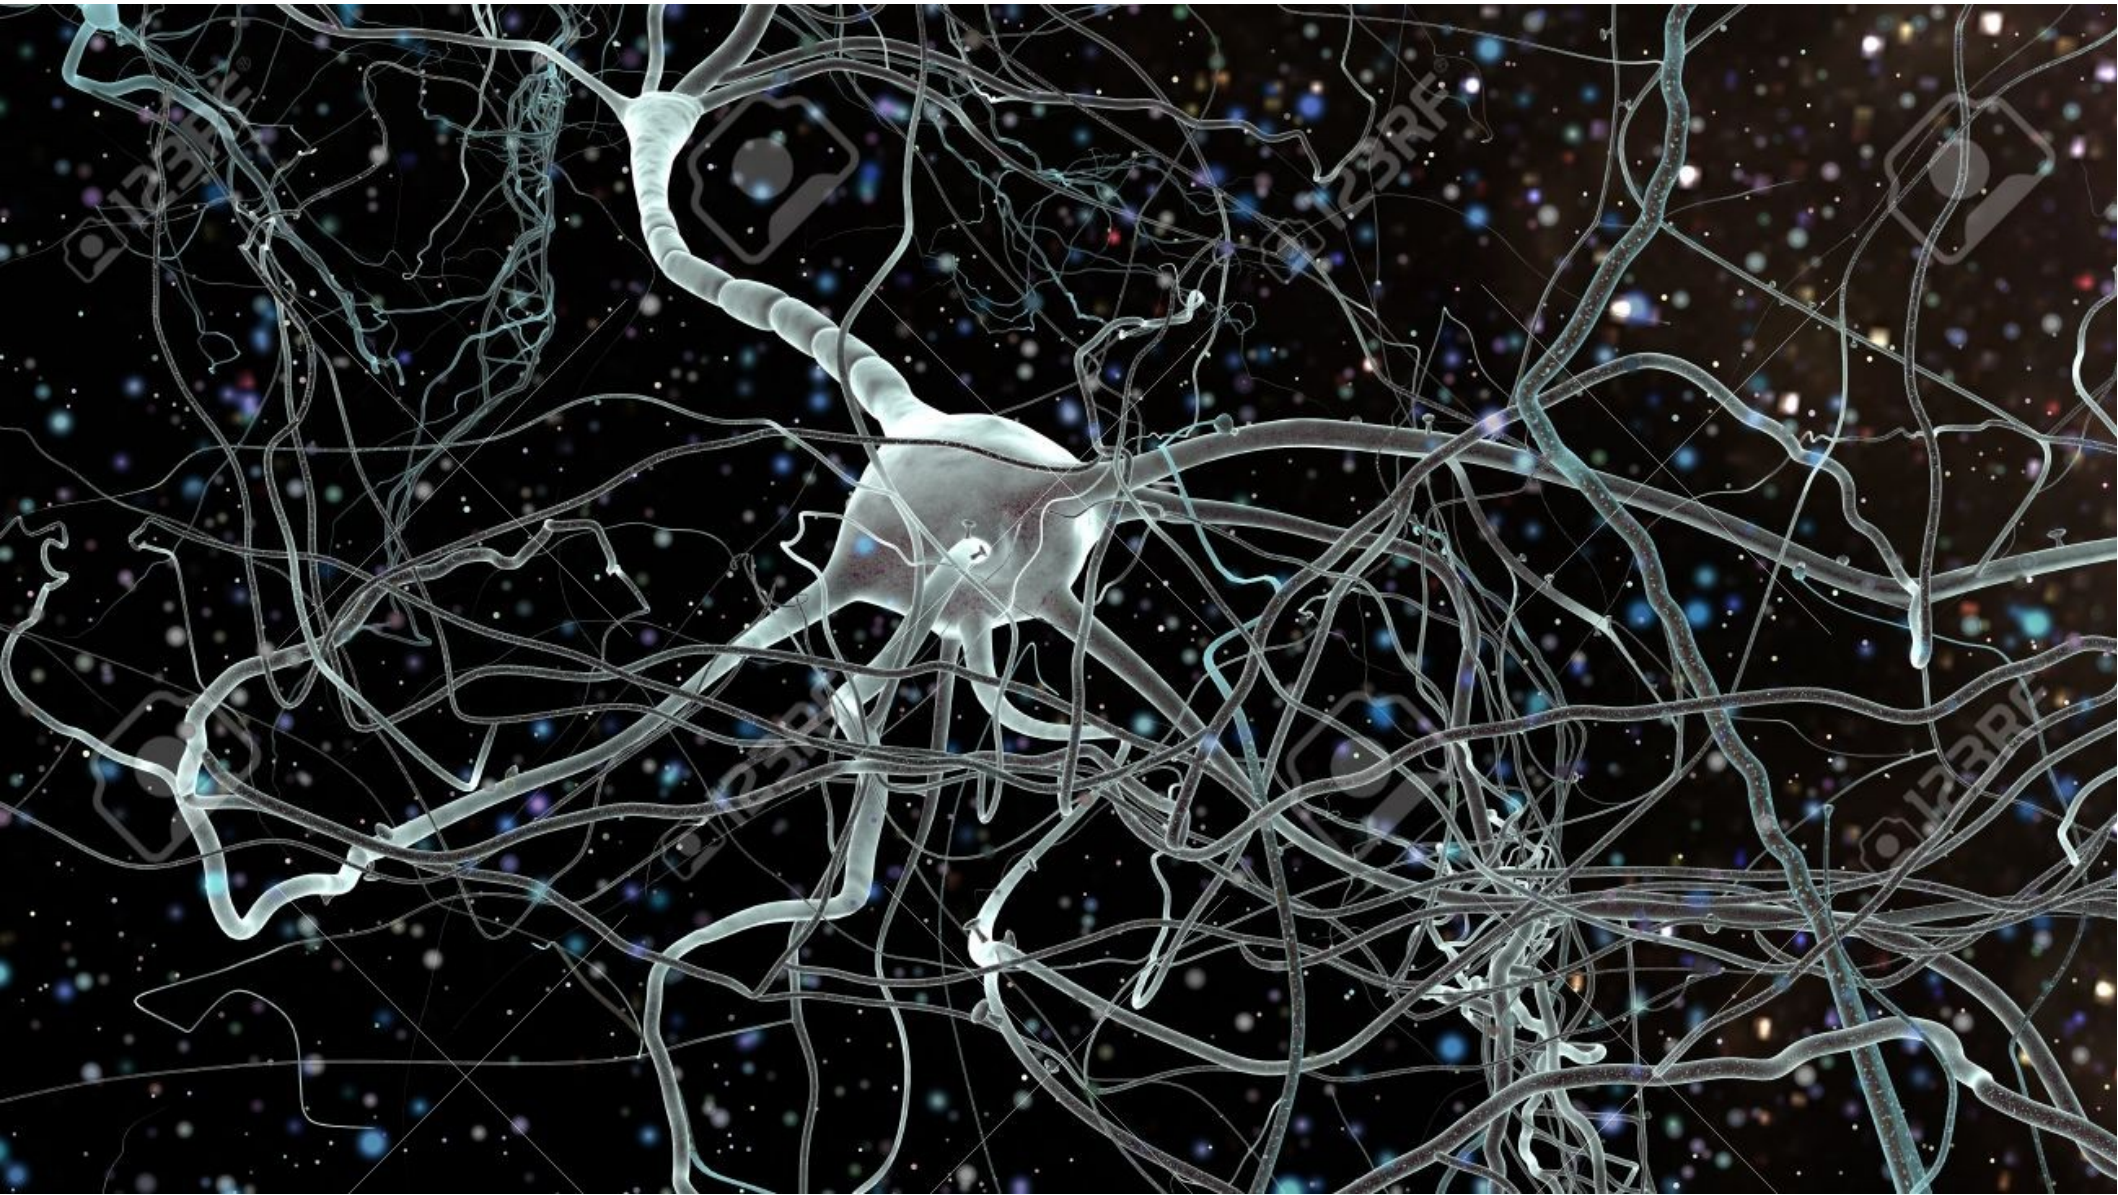
# Initialization, symmetry problem



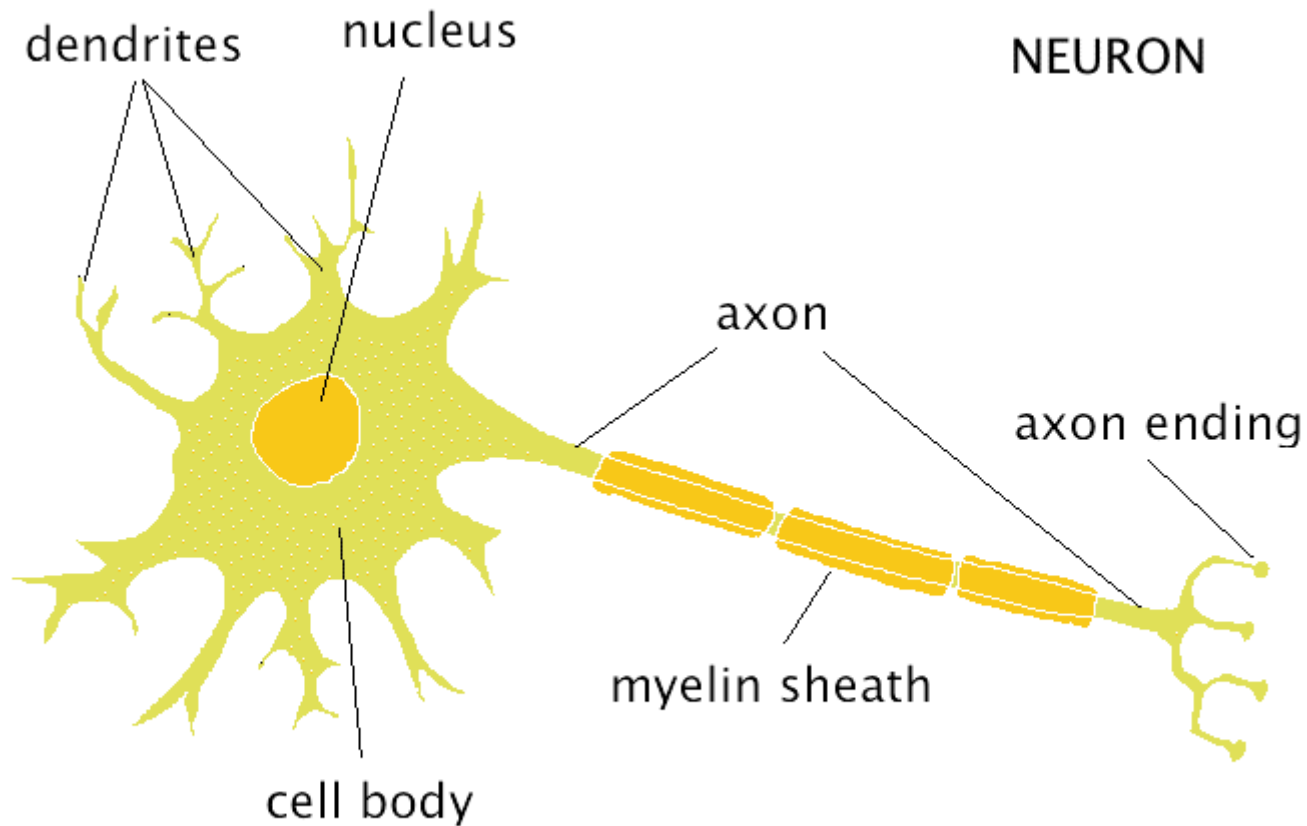
- Break the symmetry!
- Initialize with random numbers!  
 $W \leftarrow N(0, 0.01)?$   
 $W \leftarrow U(0, 0.1)?$
- Can get a bit better for deep NNs



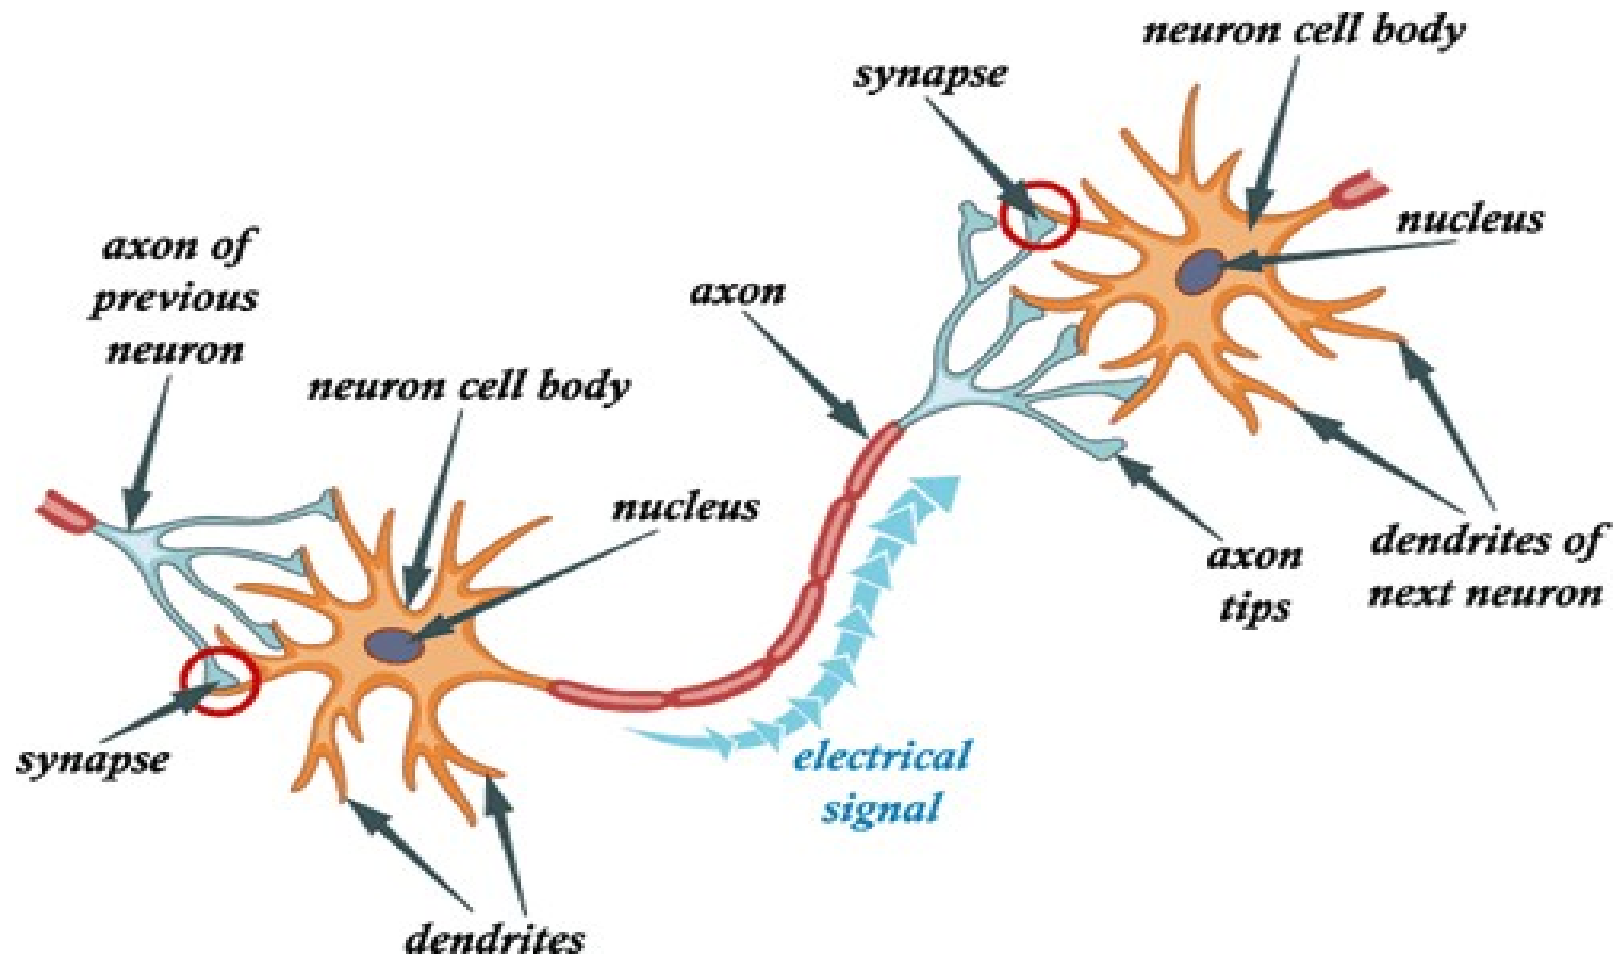
# Biological inspiration



# Biological inspiration



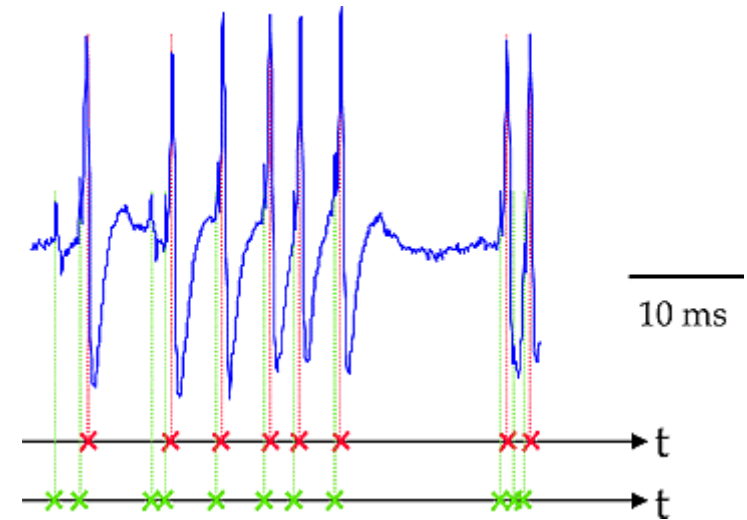
# Biological inspiration



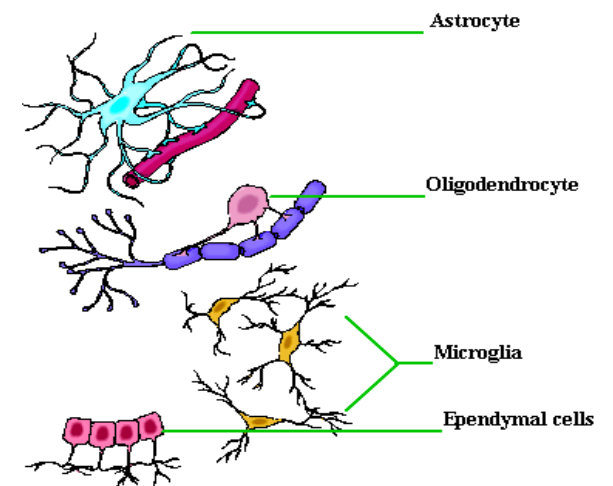


# Not actual neurons :)

- Neurons react in “spikes”, not real numbers
- Neurons maintain/change their states over time
- No one knows for sure how they “train”
- Neuroglial cells are important  
But no one knows, why



Neuroglial Cells of the CNS

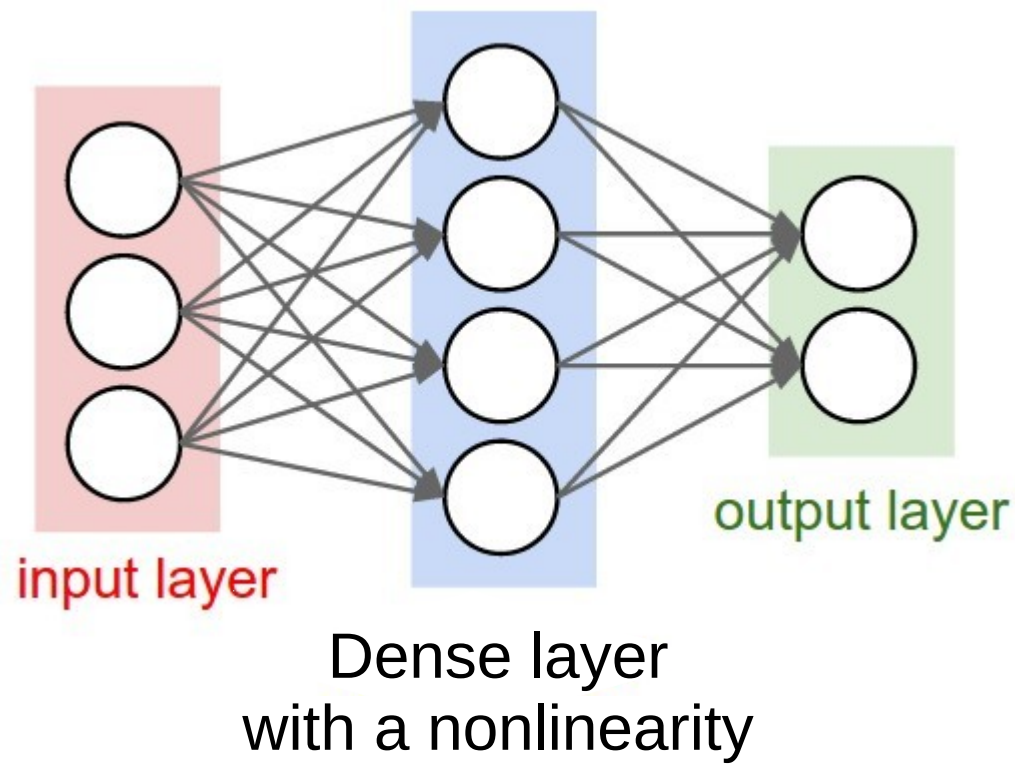


# Connectionist phrasebook

- Layer – a building block for NNs :
  - “Dense layer”:  $f(x) = Wx + b$
  - “Nonlinearity layer”:  $f(x) = \sigma(x)$
  - Input layer, output layer
  - A few more we gonna cover later
- Activation – layer output
  - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

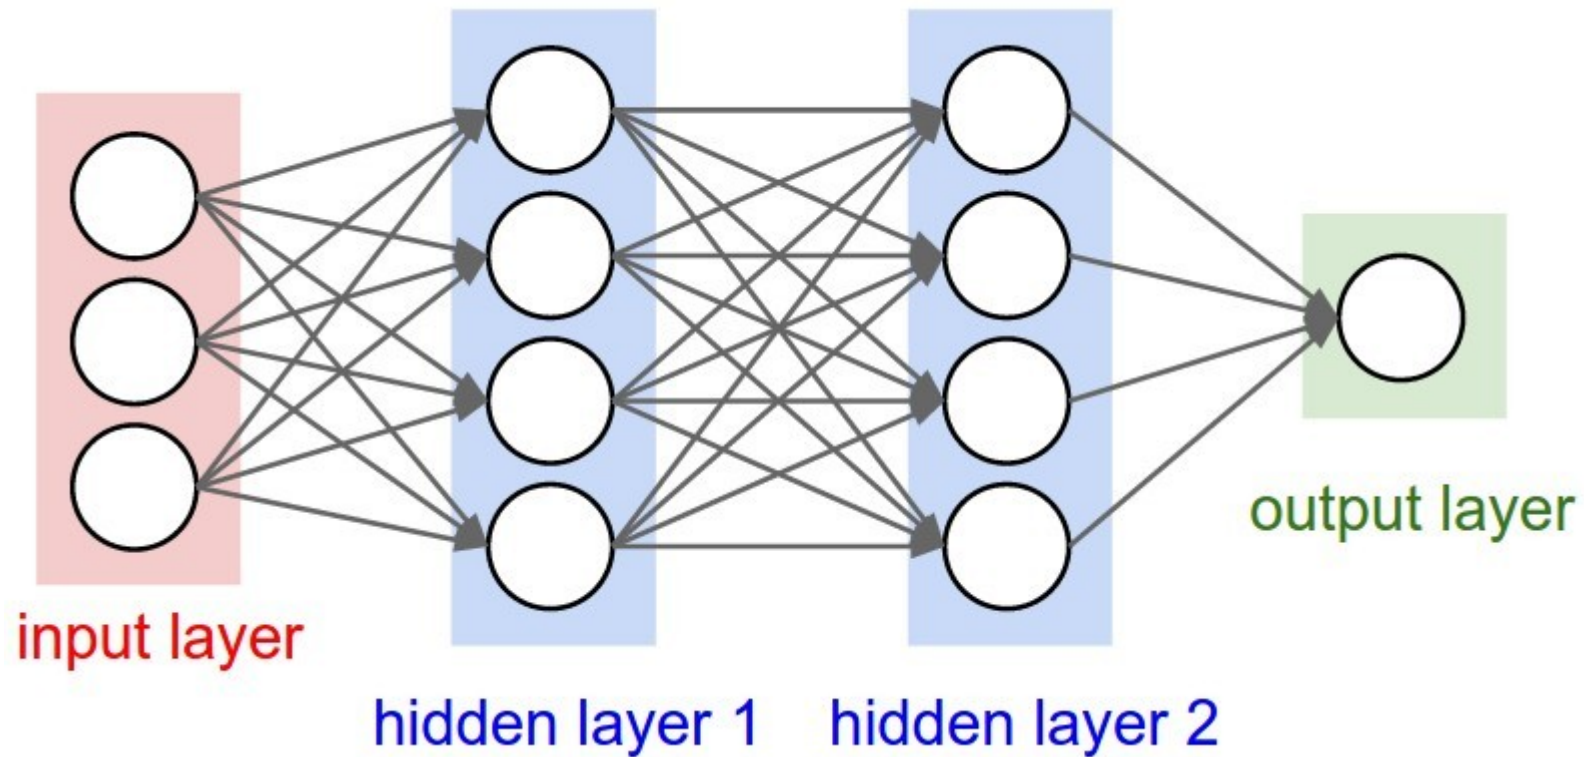


# Connectionist phrasebook



- “Train it via backprop!”

# Connectionist phrasebook



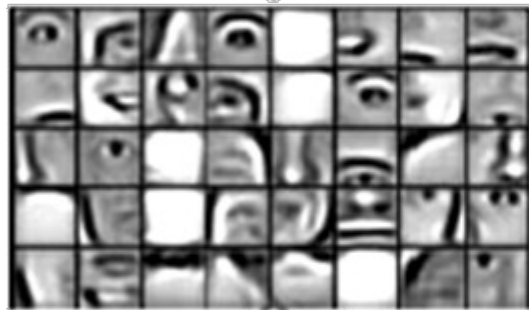
How do we train it?

# Potential caveats?

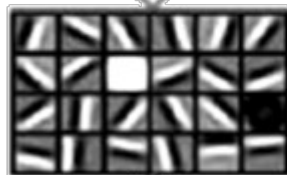


**Discrete Choices**

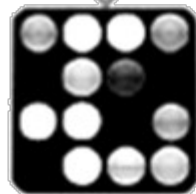
⋮



**Layer 2 Features**



**Layer 1 Features**



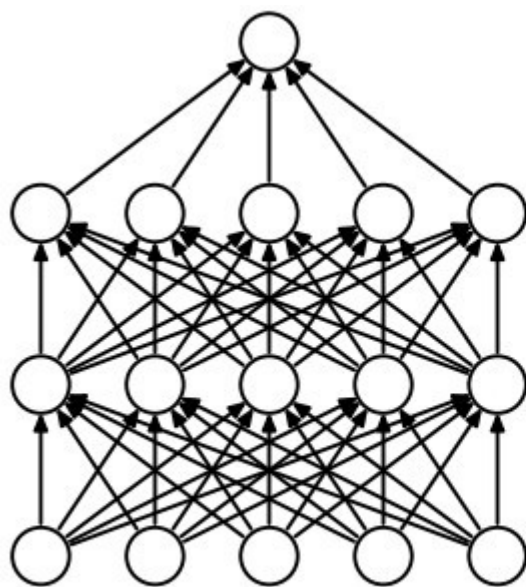
**Original Data**

# Potential caveats?

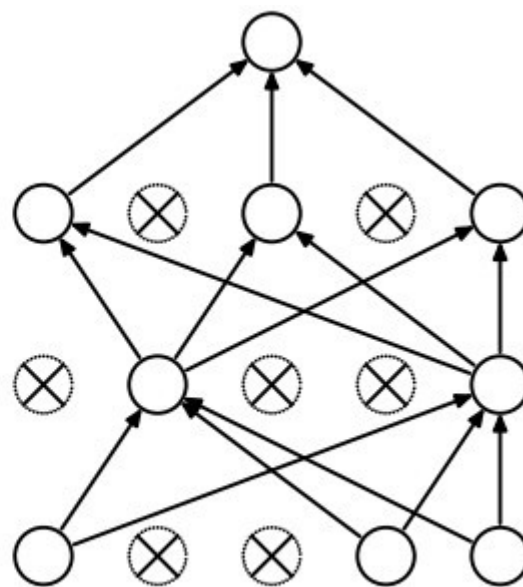
- Hardcore overfitting
- No “golden standard” for architecture
- Computationally heavy

# Regularization

- L1, L2, as usual
- Dropout



(a) Standard Neural Net



(b) After applying dropout.

# Computation

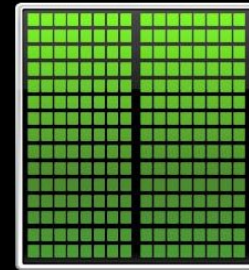


*The Difference between a CPU and GPU*



**CPU**

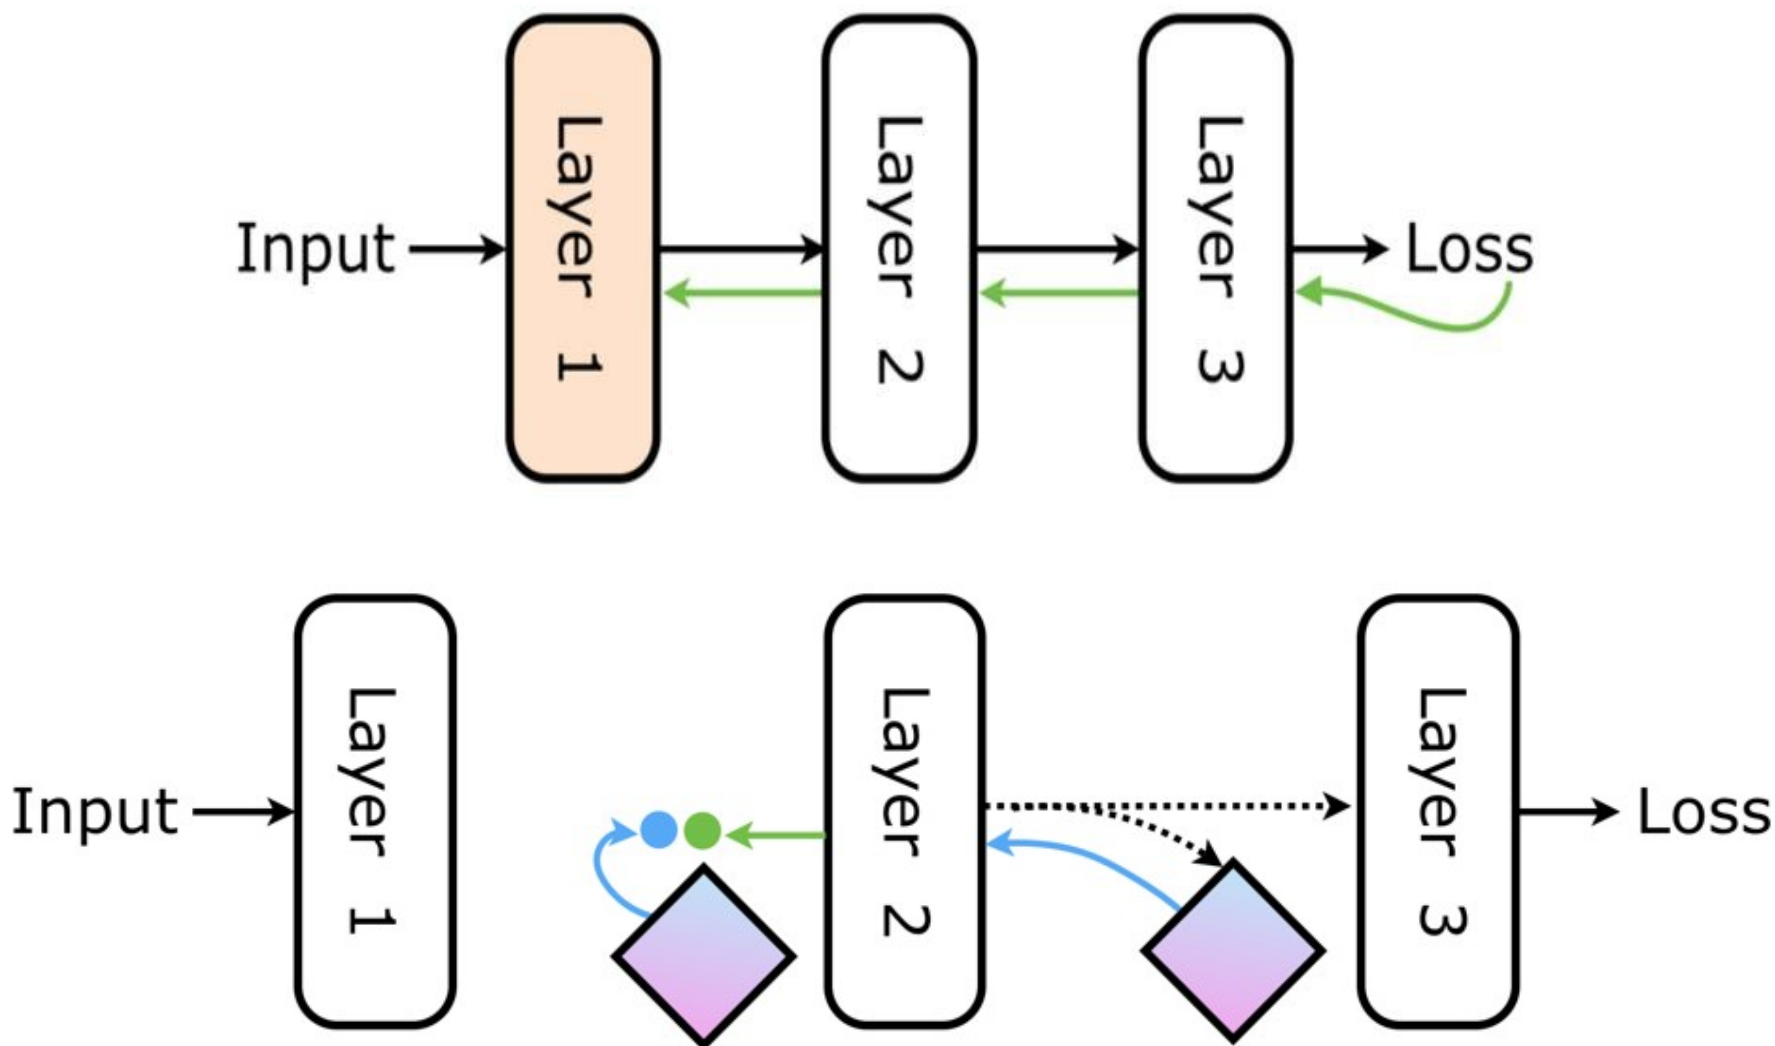
MULTIPLE CORES



**GPU**

THOUSAND OF CORES

# Is backprop the only choice?





# Nuff

Let's code some neural networks!