

# Deep Learning

## Episode 8


# Generative & Unsupervised models



# Autoencoders 101

Main idea:

- Take data in some original (high-dimensional) space;
- Project data into a new space **from which it can then be accurately restored**;
- Encoder = data to hidden
- Decoder = hidden to data
- $\text{Decoder}(\text{Encoder}(x)) \sim x$




# Why do we ever need that?

- Compress data
  - $|\text{code}| \ll |\text{data}|$
- Dimensionality reduction
  - Before feeding data to your XGBoost

<to be continued>

# Matrix decompositions

- Example: SVD/PCA

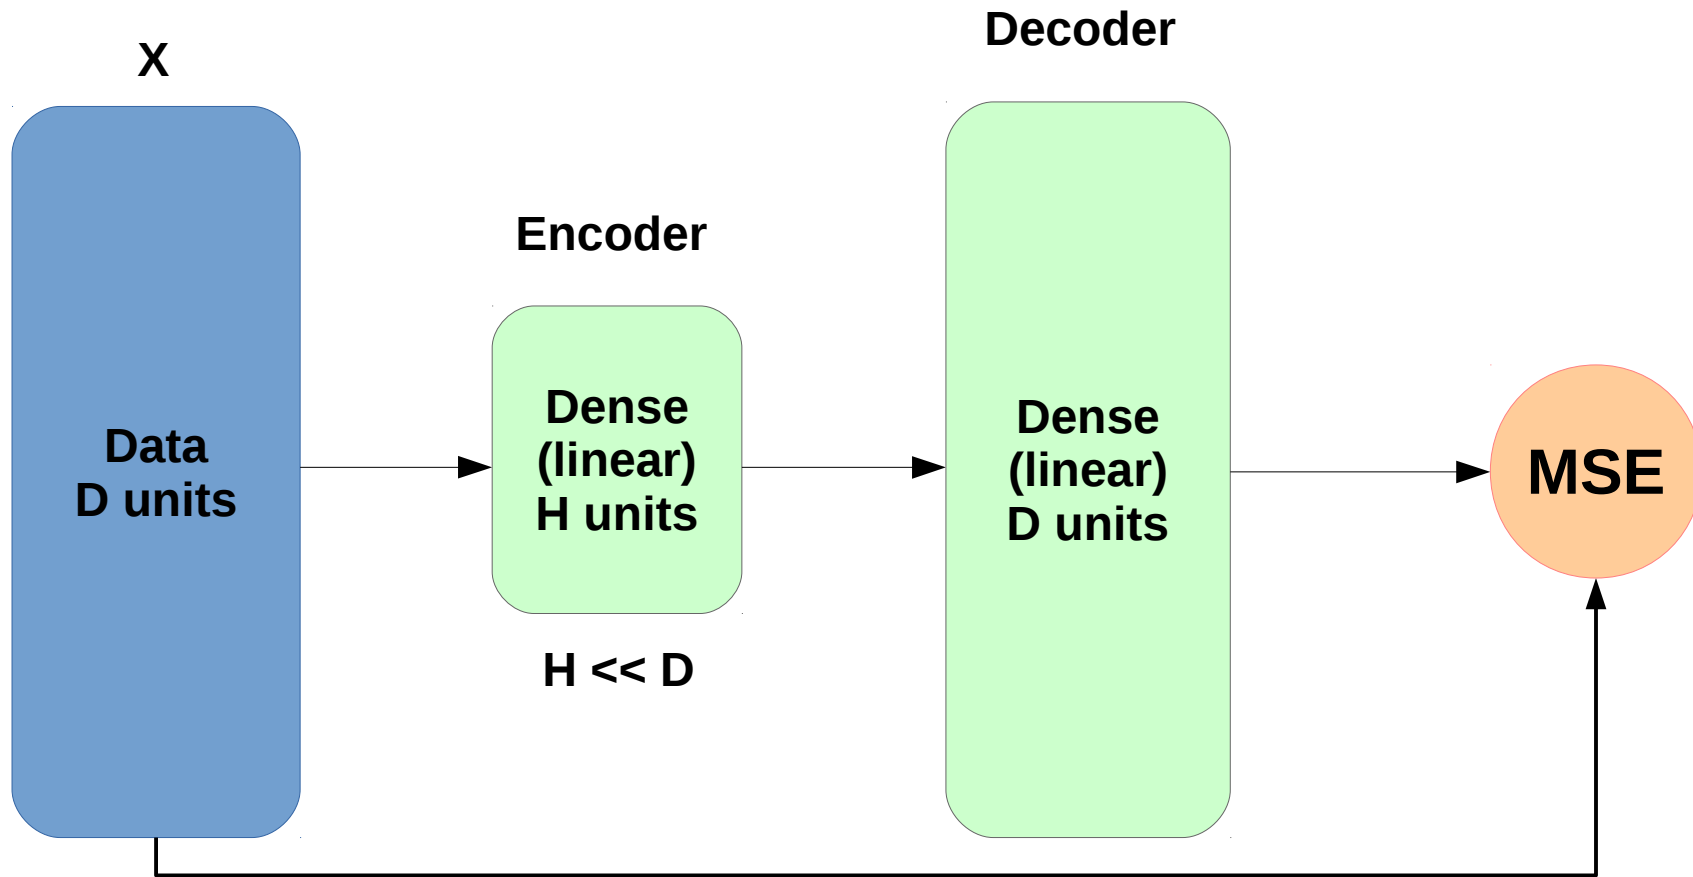



- Minimizing reconstruction error

$$L = \|X - U \cdot S \cdot V^T\|$$

# Matrix decomposition

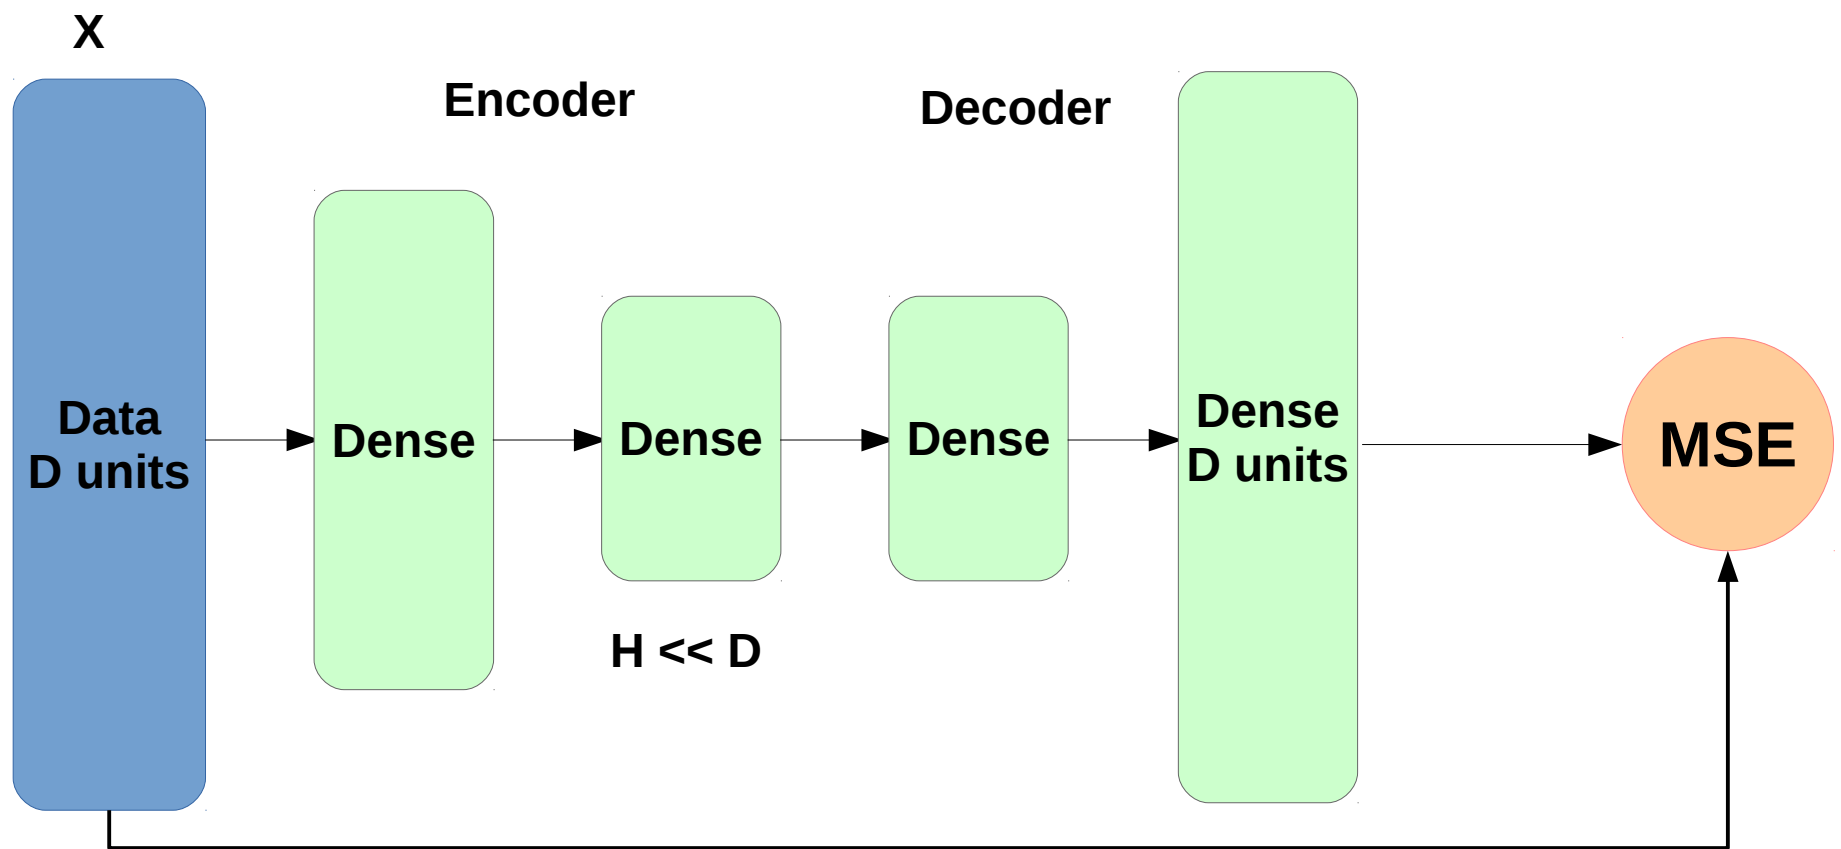
- A different perspective





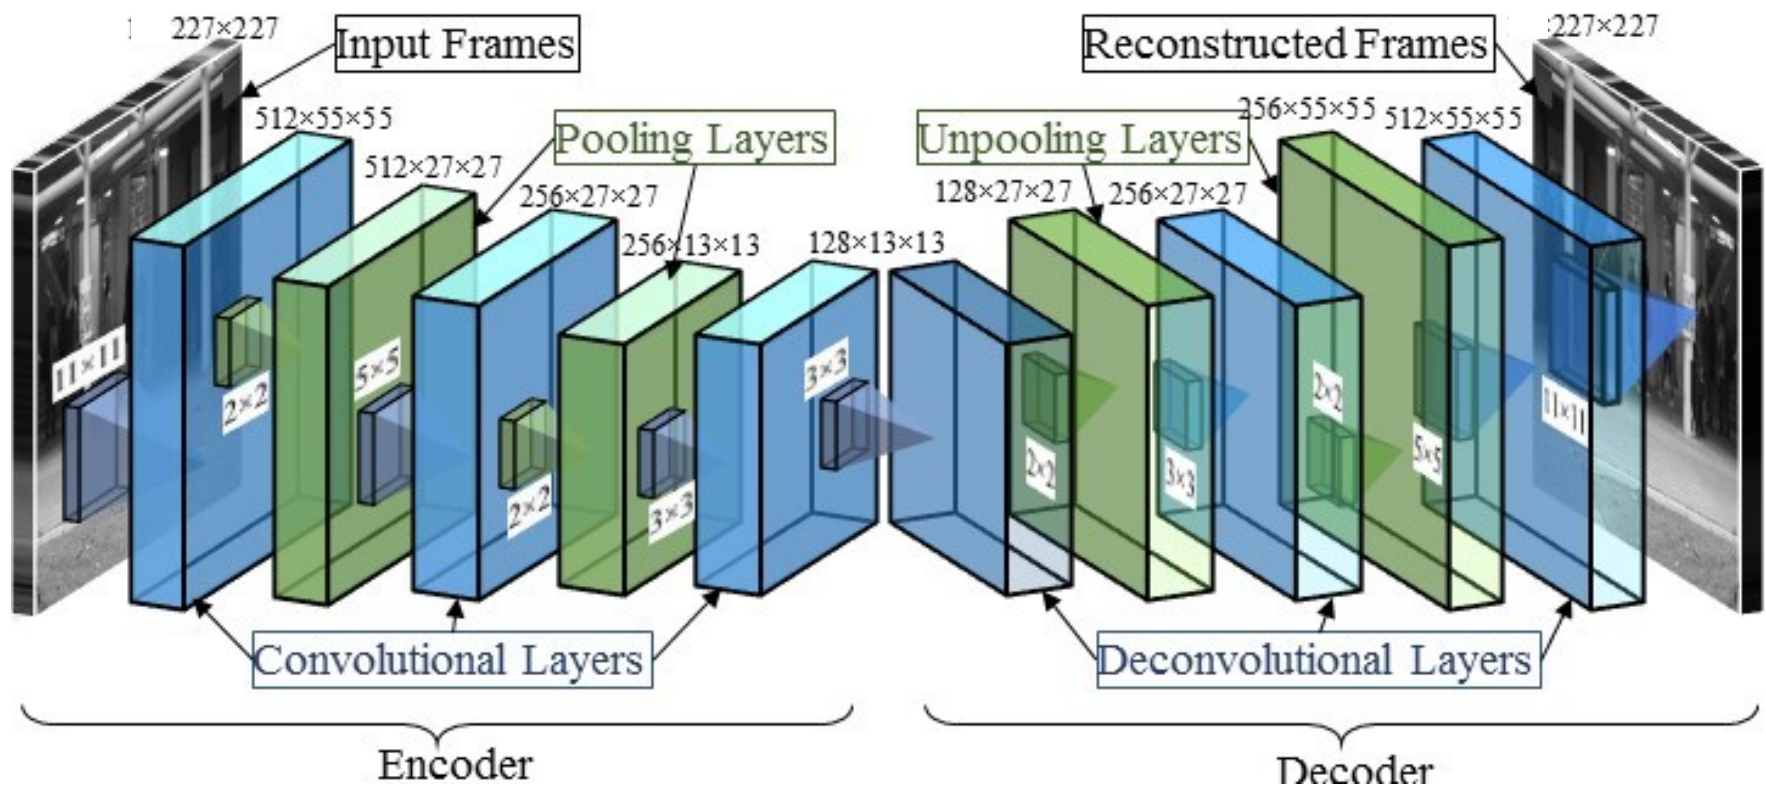
# (kinda) Deep autoencoder

- Stack more layers!



Quiz: What if data is an image?

# Image2image: fully-convolutional



**Quiz: what is the compression rate here?**

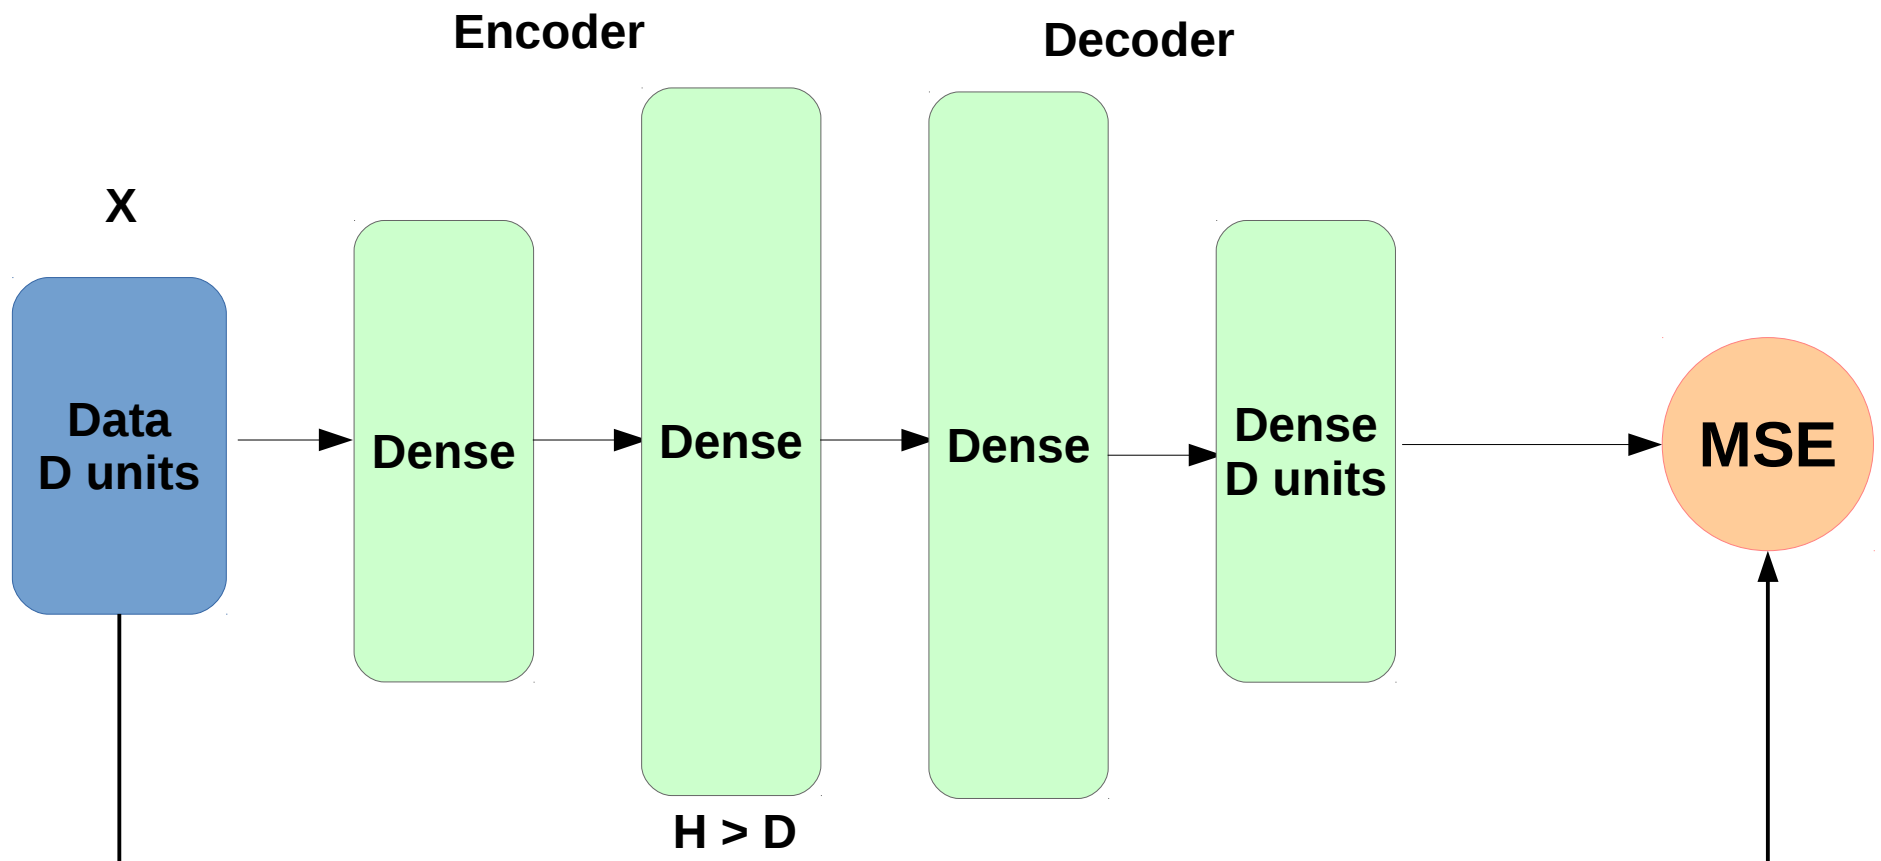


# Why do we ever need that?

- Compress data
  - $|\text{code}| \ll |\text{data}|$
- Dimensionality reduction
  - Before feeding data to your XGBoost
- **Learn some great features!**
  - Before feeding data to your XGBoost

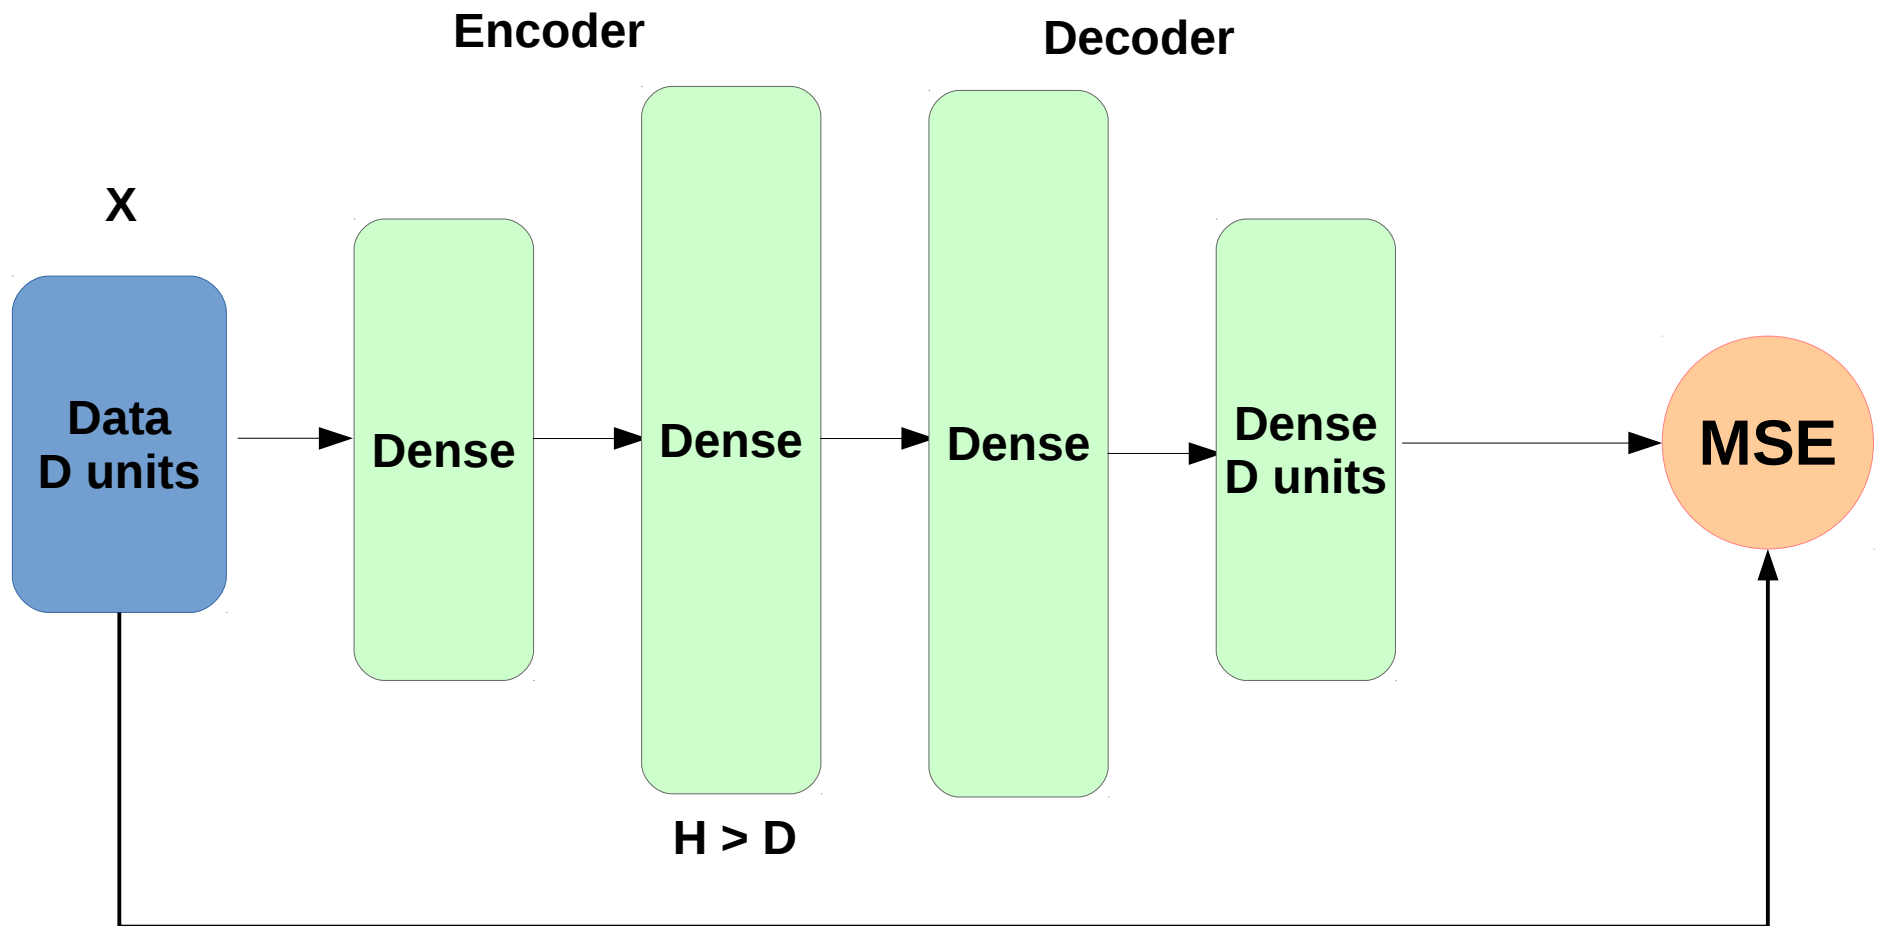
# Expanding autoencoder

- Bigger/richer representation



# Expanding autoencoder


- Bigger/richer representation



Something's wrong with this guy. **Ideas?**

# Expanding autoencoder


- Naive approach will learn identity function!
- Gotta regularize!



$$L = \|X - Dec(Enc(X))\|$$

# Sparse autoencoder


- Naive approach will learn identity function!
- Idea 1: L1 on **activations**, sparse code



$$L = ||X - Dec(Enc(X))|| + \sum_i |Enc_i(X)|$$

# Redundant autoencoder


- Naive approach will learn identity function!
- Idea 2: noise/dropout, redundant code



$$L = \|X - \text{Enc}(\text{Noise}(\text{Dec}(X)))\|$$

# Denoizing autoencoder


- Naive approach will learn identity function!
- Idea 3: distort input, learn to undo distortion




$$L = \|X - \text{Enc}(\text{Dec}(\text{Noise}(X)))\|$$

# Sparse Vs Denoizing

- Filter weights, 12x12 patches



Sparse AE



Denoizing AE

Actually meaningless :)

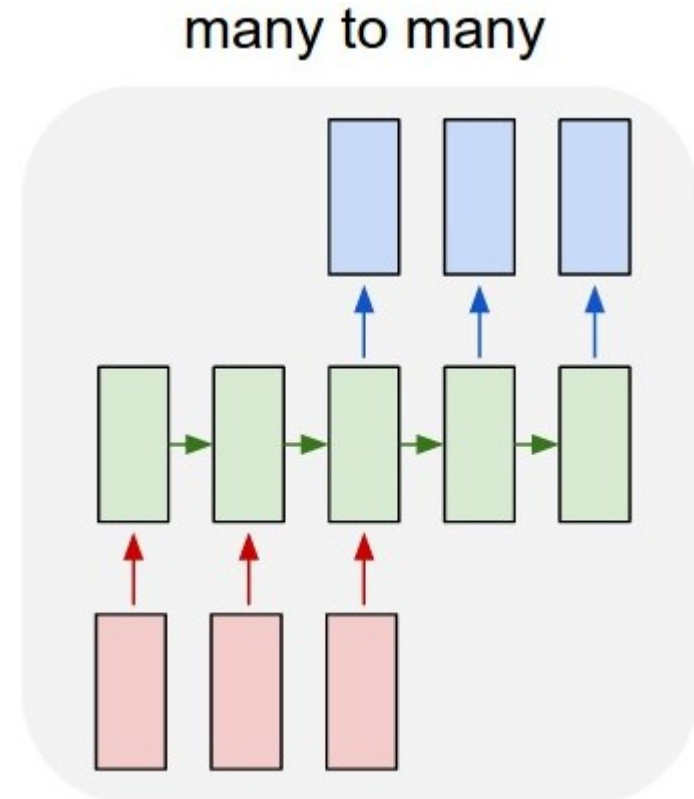


# Why do we ever need that?

- Compress data
  - $|\text{code}| \ll |\text{data}|$
- Dimensionality reduction
  - Before feeding data to your XGBoost
- Learn some great features!
- **Unsupervised pretraining**
  - Large amounts of data
  - Features may be irrelevant

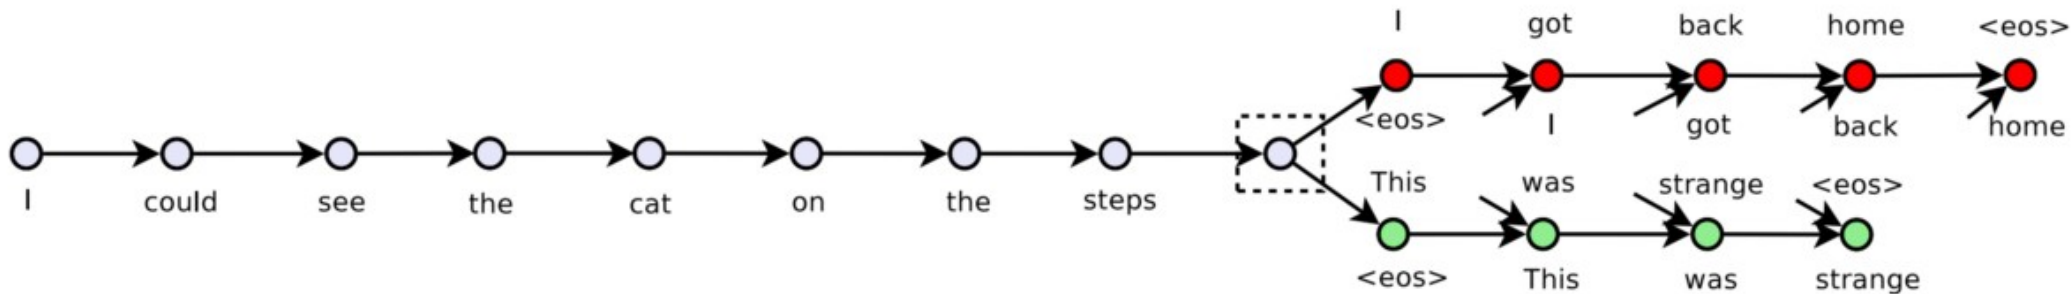
# Recurrent autoencoders

- Regular encoder-decoder
- Where is the bottleneck?
- How do we train it?



# Skip-thought

- Word2vec **skip-gram**:
  - Word → neighboring words
  - Embedding + Dense
  - Word vectors
- Phrase2vec **skip-thought**:
  - Sentence → prev/next sentence
  - Encoder-decoder
  - Sentence vectors




# Why do we ever need that?

- Compress data
  - $|\text{code}| \ll |\text{data}|$
- Dimensionality reduction
  - Before feeding data to your XGBoost
- Learn some great features!
- Unsupervised pretraining
  - Large amounts of data
  - Features may be irrelevant
- **Generating new images!**

# Image morphing with AE

Idea:

- If  $\text{Enc}(\text{image1}) = c1$   
 $\text{Enc}(\text{image2}) = c2$
- Than maybe  $(c1+c2)/2$  is a semantic average of the two images



# Image morphing with AE

Idea:

- Look for a common direction vector for “add mustache” or “add age” changes.
- Apply to new images



+ OLD =



- FEMALE  
+ MALE =




# Mean Squared Error

Pixelwise MSE:

- A “**cat on the left**” is closer to “**dog on the left**” than to “**cat on the right**”
- We may want to avoid that effect
- Can we obtain image representation that is less sensitive to small shifts?

# Sketch: using pre-trained nets




$$L = \|f(X) - f(Dec(Enc(X)))\|$$




# Image generation

- Chairs (type, view, orientation)




- Classifier



# Image generation

- Generator



**Problem:** MSE sucks at this task.

**Ideas?**

**WHAT IF WE TRAIN  
THAT 2-ND NETWORK**



**TO HELP US TRAIN  
THE FIRST NETWORK**

# Generative Adversarial Networks

## Generator




Generate image  
(should be plausible)

## Discriminator




Tell if image is plausible  
(image)  $\rightarrow$   $P(\text{fake})$



# Generative Adversarial Networks

- Generator




- Discriminator



# Generative Adversarial Networks

- Generator




- Discriminator

$$L_D = -\log[1 - \text{Disc}(\text{real data})] - \log \text{Disc}(\text{Gen}(\text{seed}))$$



# Generative Adversarial Networks

- Generator  $L_G = -\log[1 - \text{Disc}(\text{Gen}(\text{seed}))]$



- Discriminator

$$L_D = -\log[1 - \text{Disc}(\text{real data})] - \log \text{Disc}(\text{Gen}(\text{seed}))$$





# Generative Adversarial Networks

for number of training iterations do

for  $k$  steps do

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log \left( 1 - D(G(z^{(i)})) \right) \right].$$

end for

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(z^{(i)})) \right).$$

end for

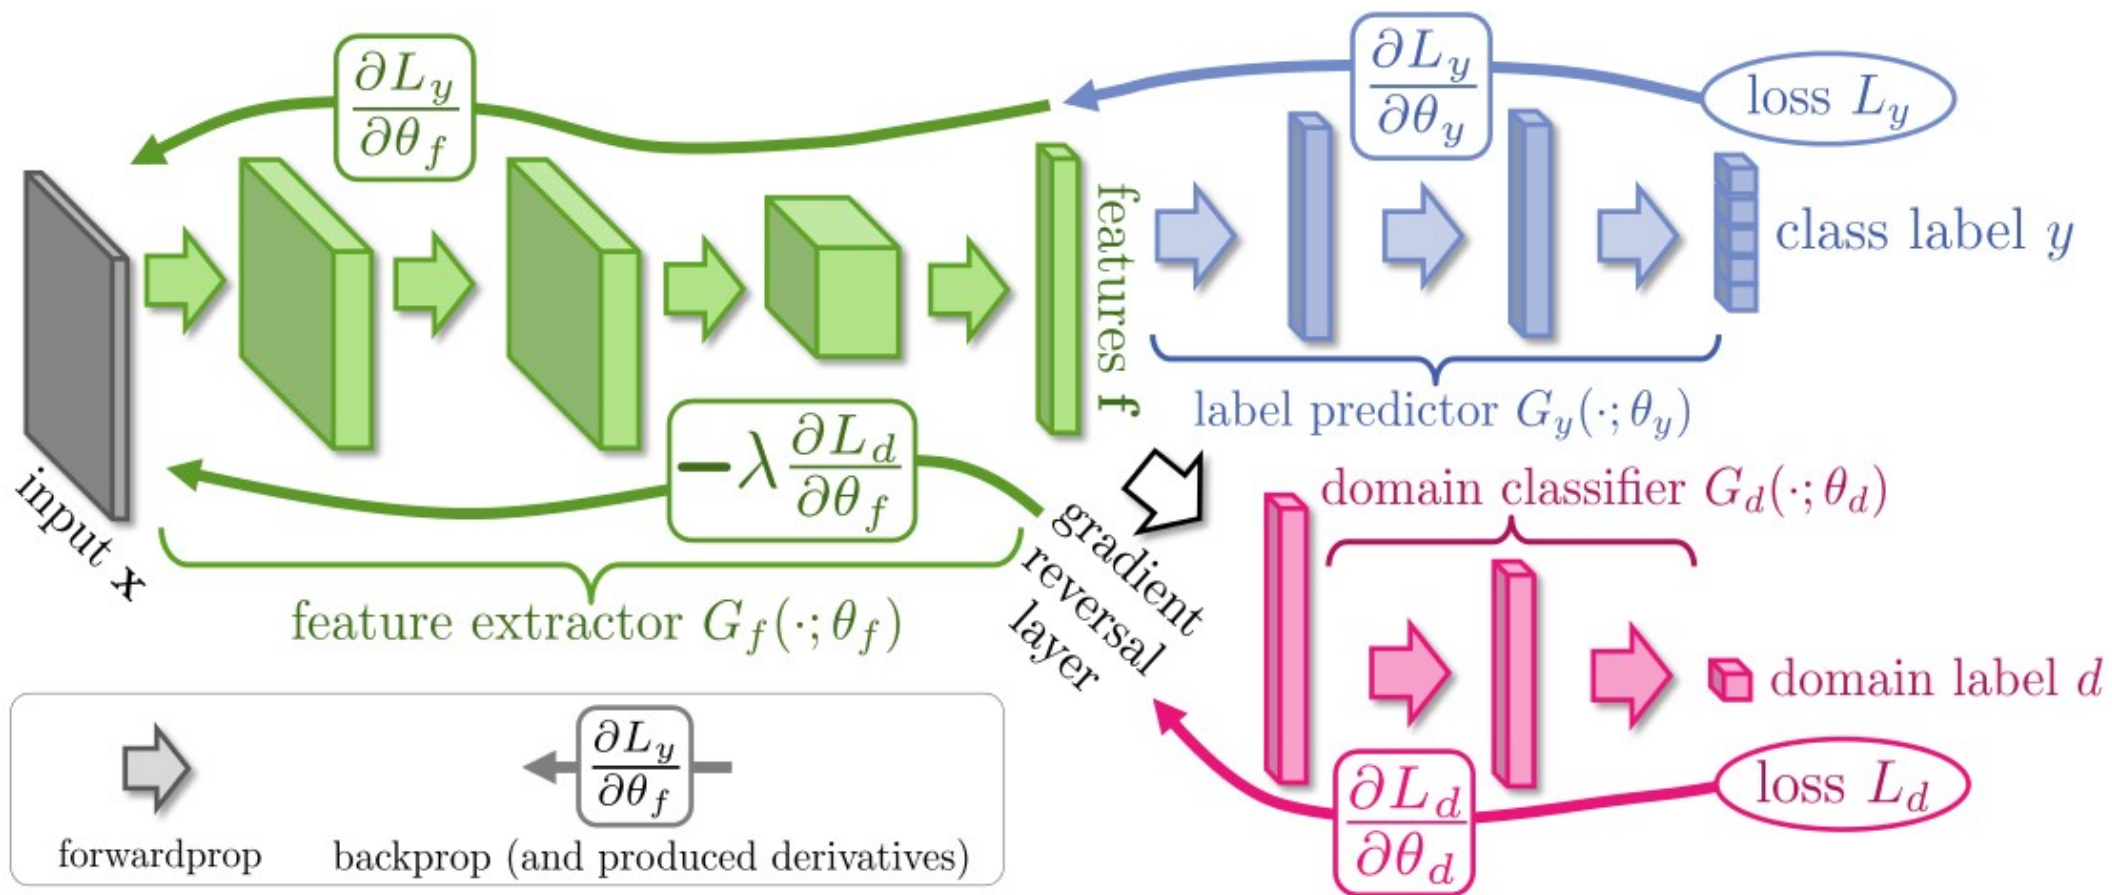
# Adversarial domain adaptation

- Two domains
  - e.g. mnist digits Vs actual digits on photos
- First domain is labeled, second is not
- Wanna learn for the second domain

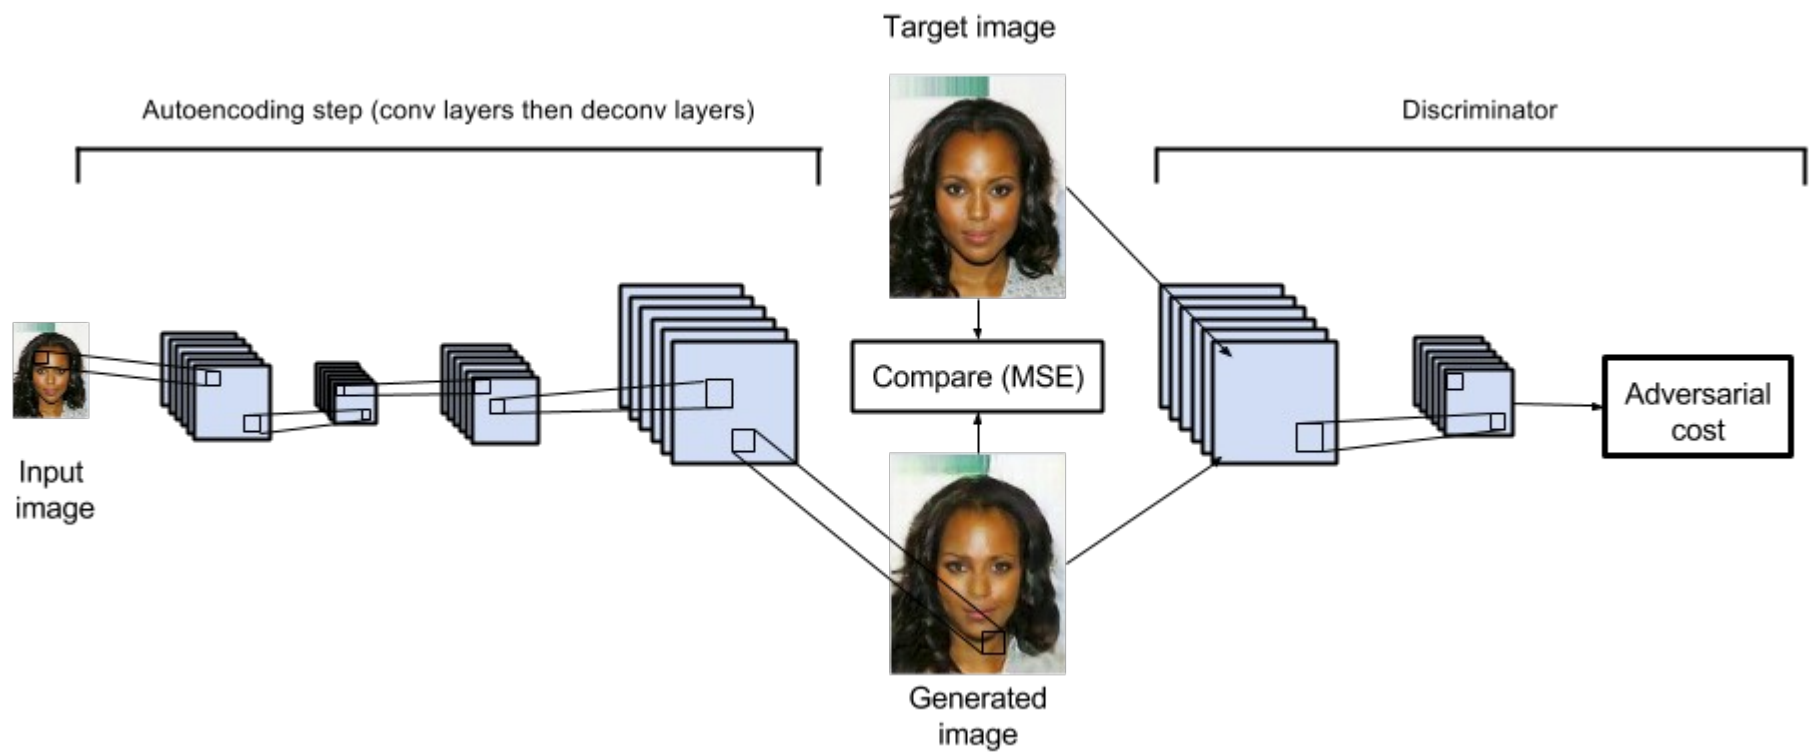


# Domain adaptation

- Idea: discriminator should not be able to distinguish features on two domains



# Adversarial autoencoders



# Brace yourselves



# Art style transfer

- Ideas?

