

Reinforcement learning

Episode 2

Approximate & deep RL




Yandex
Data Factory

LAMBDA 



**British Hedgehog
Preservation Society**

Recap: MDP



Classic MDP(Markov Decision Process)

Agent interacts with environment

- Environment states: $s \in S$
- Agent actions: $a \in A$
- State transition: $P(s_{t+1}|s_t, a_t)$
- Reward: $r_t = r(s_t, a_t)$

Recap: total reward

Objective:
Total reward

$$R_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots + \gamma^n \cdot r_{t+n}$$

$\gamma \in (0, 1)$ const
 $\gamma \sim$ patience

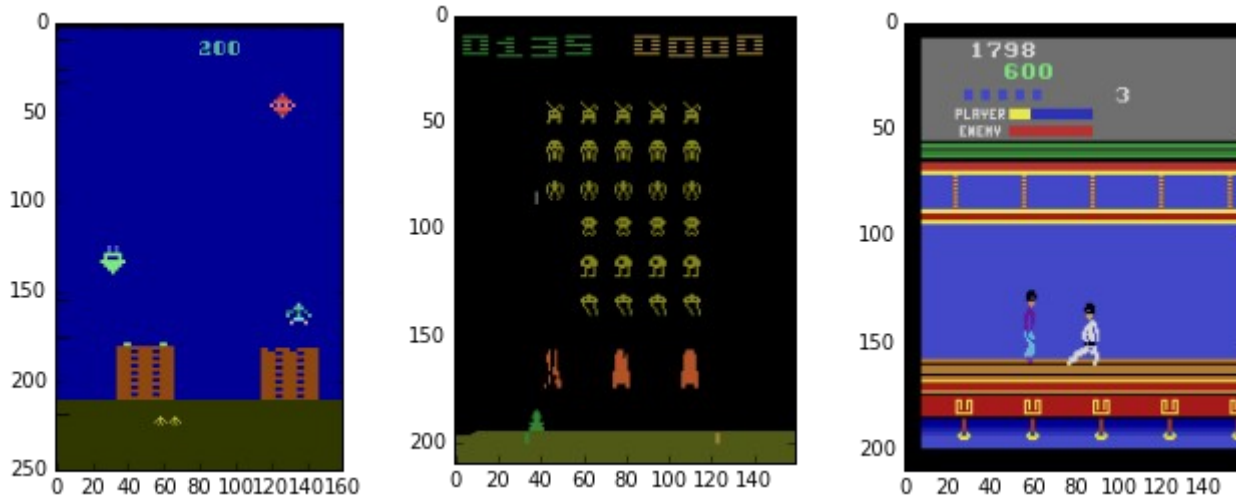
Reinforcement learning:

- Find policy that maximizes the expected reward

$$\pi = P(a|s) : E[R] \rightarrow \max$$




Reality check: videogames



- **Trivia:** What are the states and actions?


Recap: Q-learning



Definition: $Q(s, a)$ is an expected total reward R that can be obtained starting from state s by taking action a and following optimal policy since next state.

$$\pi(s) : \operatorname{argmax}_a Q(s, a)$$

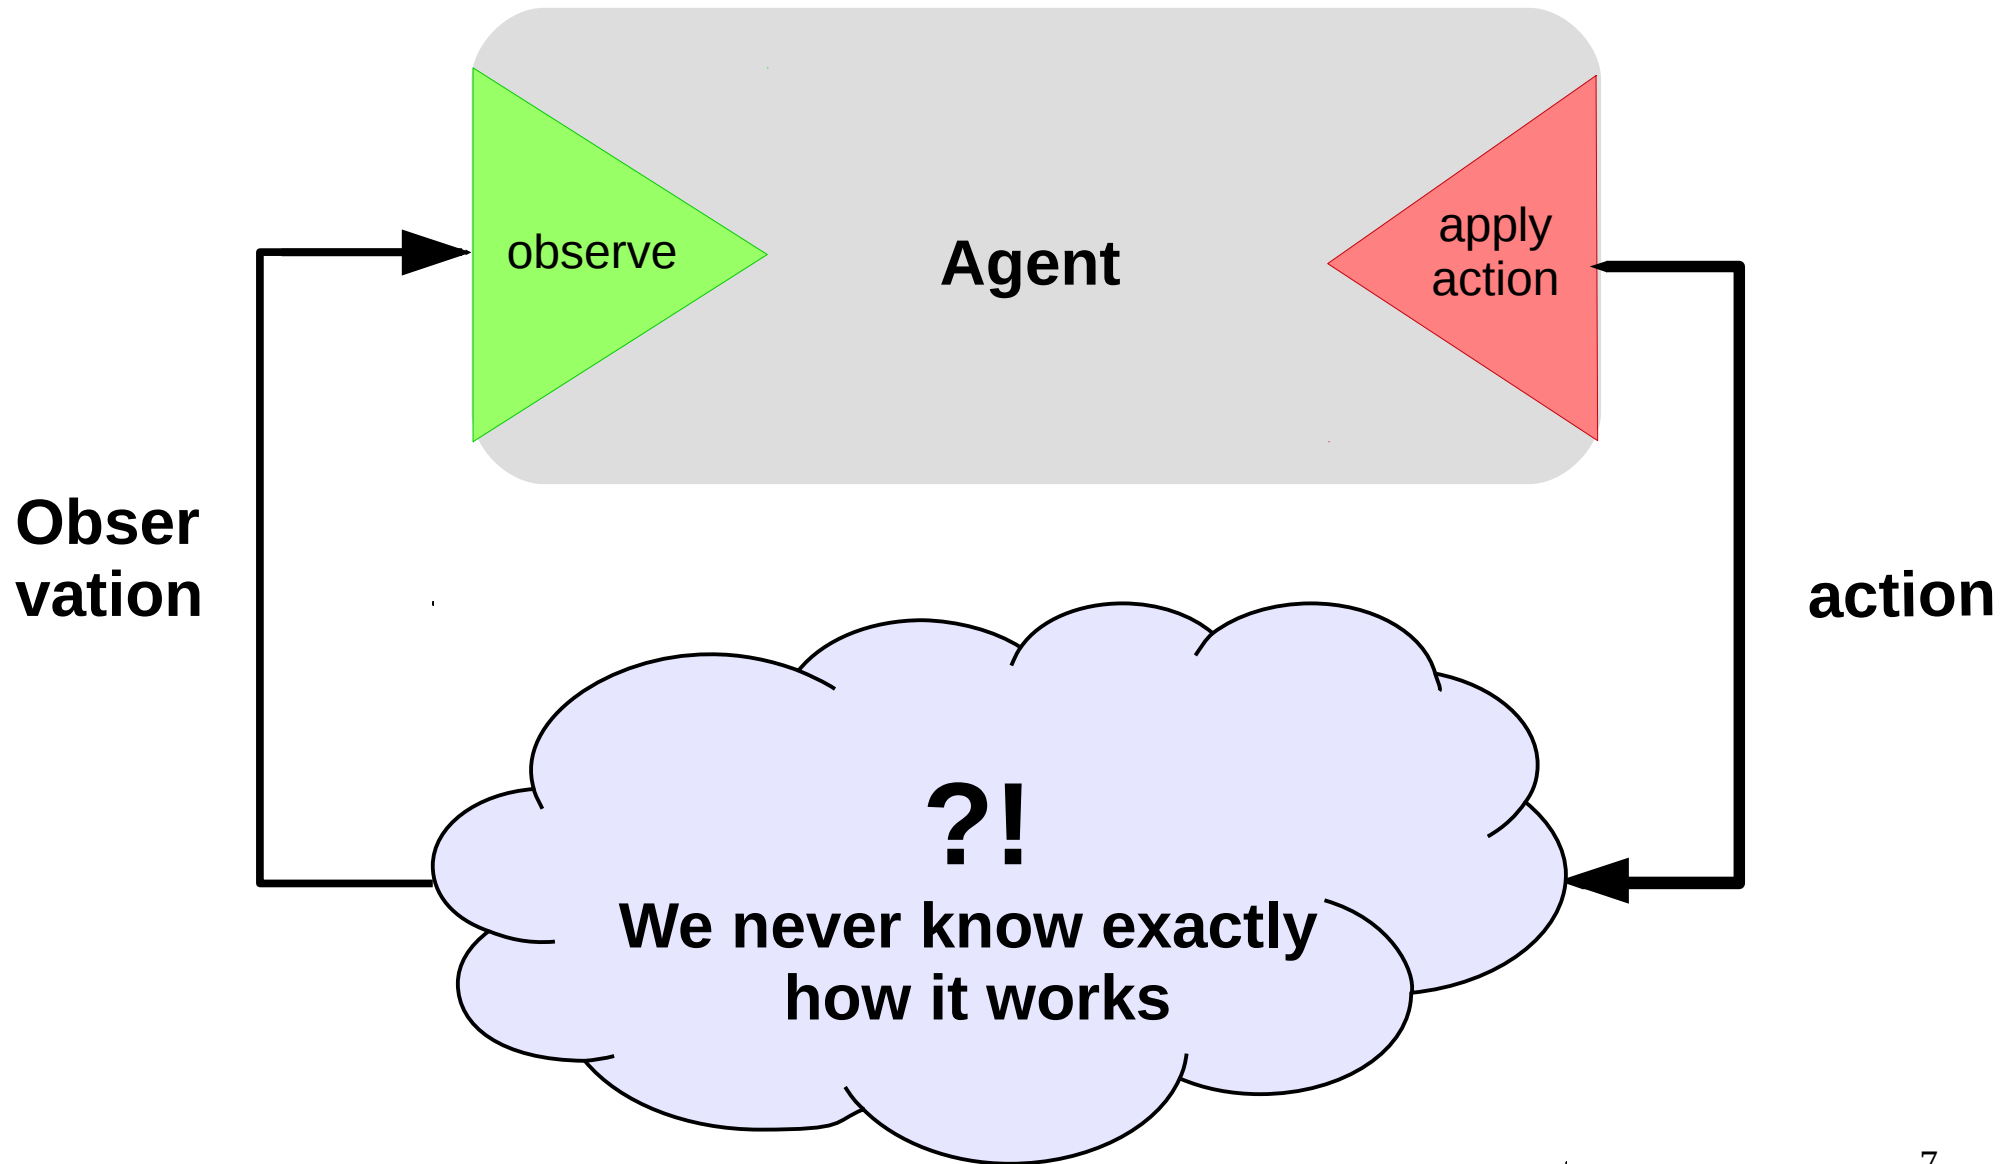
Recap: Q-learning



- Initialize with zeros/random
- Iteratively minimize

$$\operatorname{argmin}_{Q(s_t, a_t)} \left(Q(s_t, a_t) - [r_t + \gamma \cdot \max_{a'} Q(s_{t+1}, a')] \right)^2$$

Real world



Problem:

State space is usually large,
sometimes continuous.

And so is action space;

However, states do have a structure, similar
states have similar action outcomes.

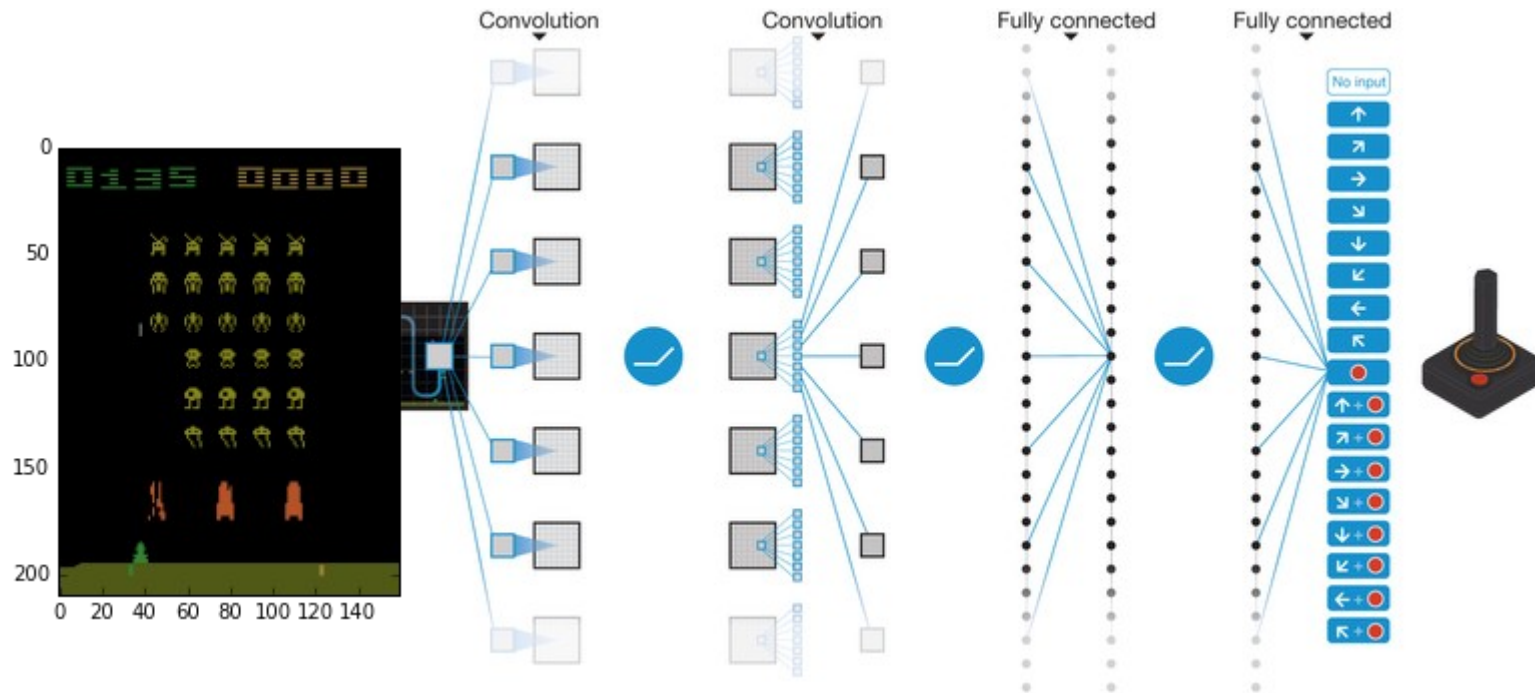
From tables to approximations

- Before:
 - For all states, for all actions, remember $Q(s,a)$
- Now:
 - Approximate $Q(s,a)$ with some function
 - e.g. linear model over state features

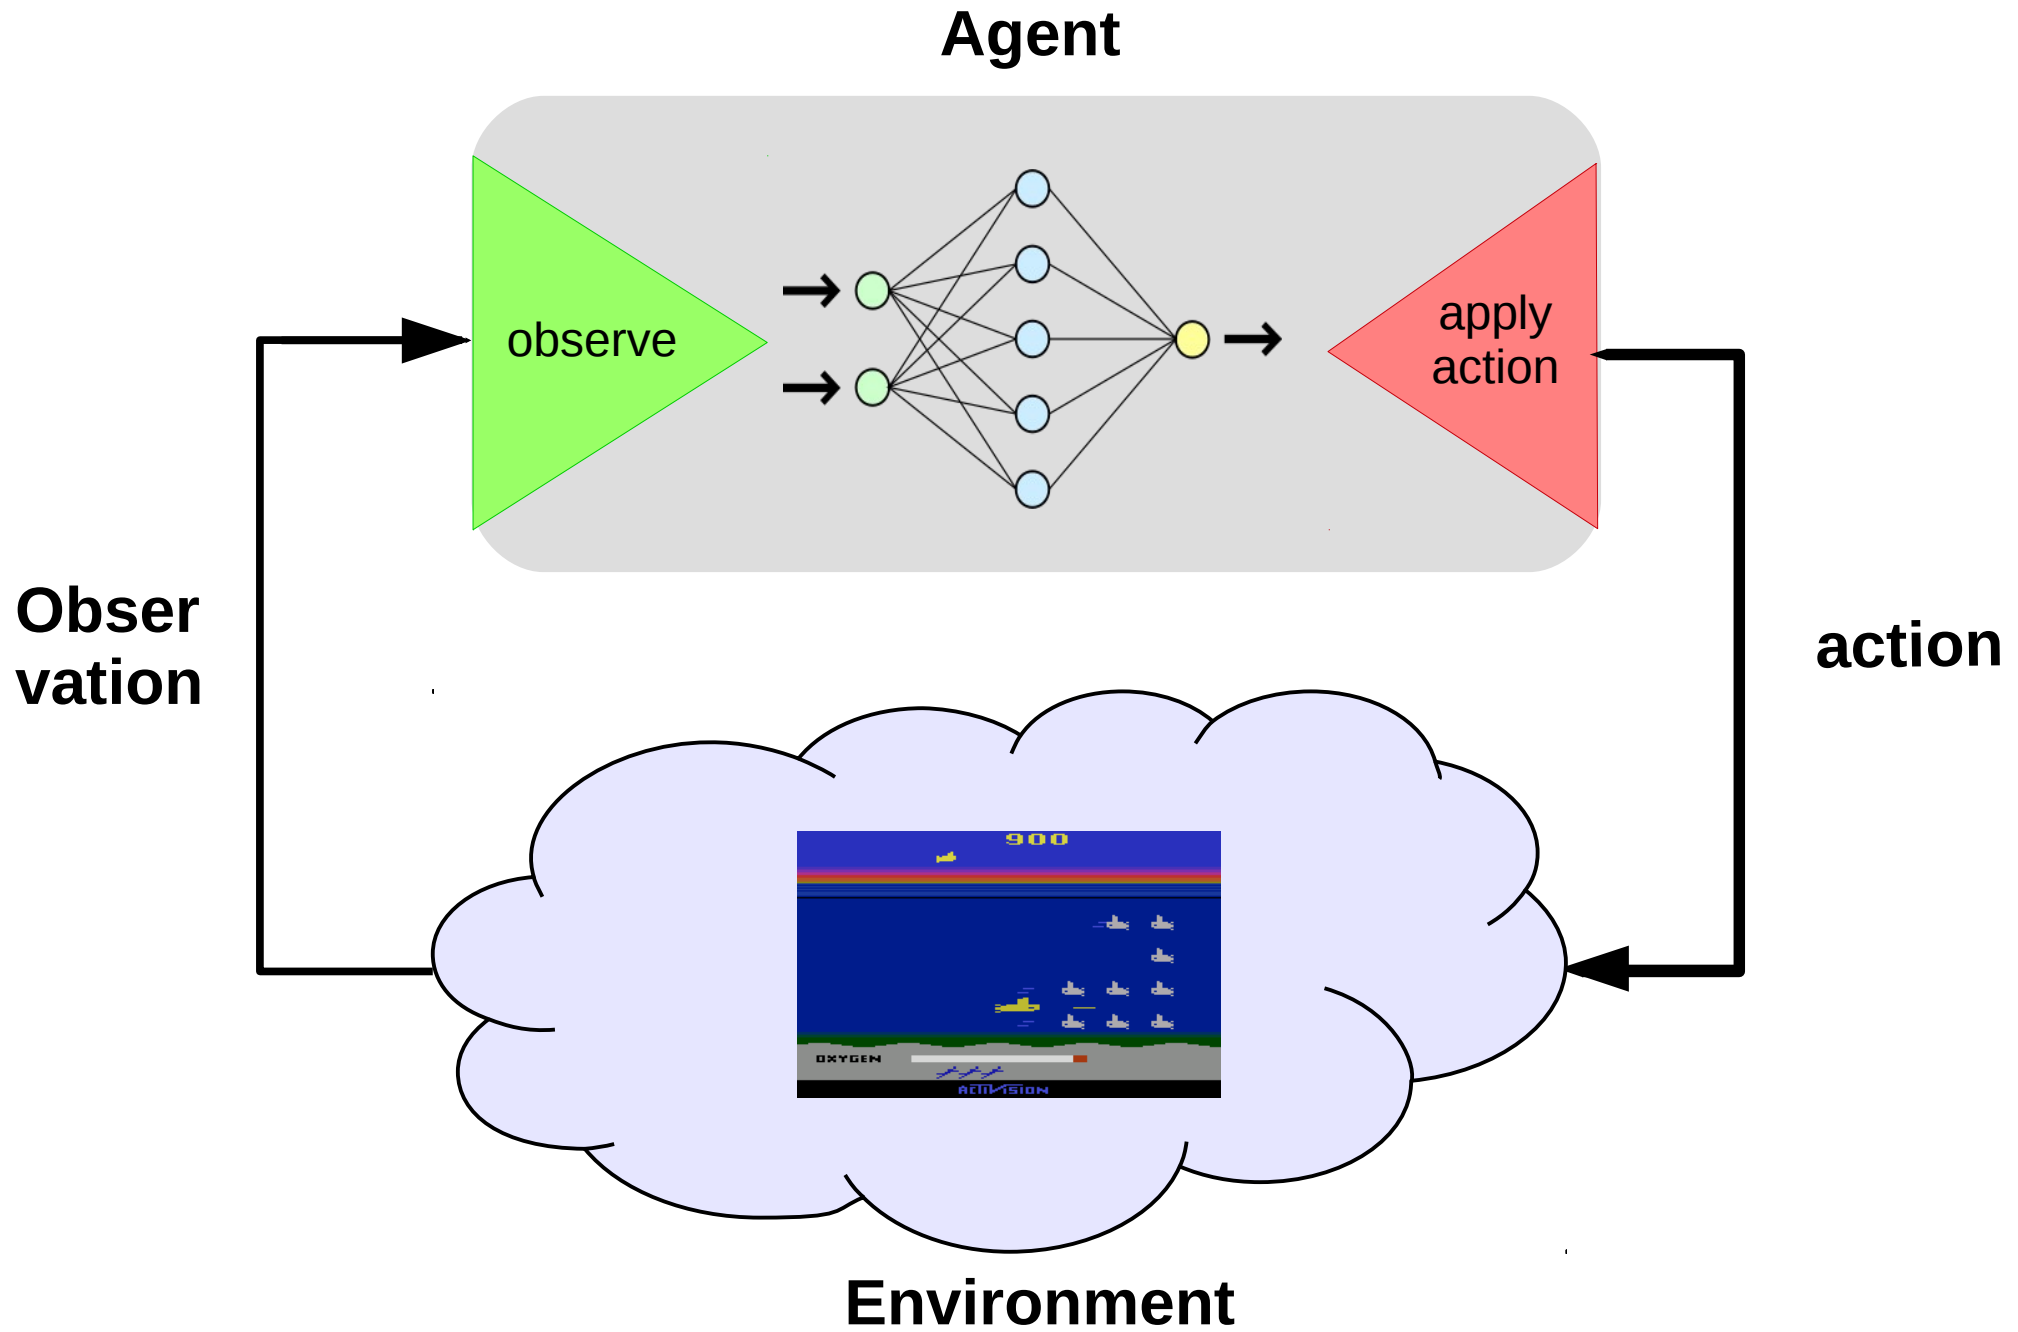
$$\operatorname{argmin}_{w,b} \left(Q(s_t, a_t) - [r_t + \gamma \cdot \max_{a'} Q(s_{t+1}, a')] \right)^2$$

Trivia: should we use linear regression or logistic regression?

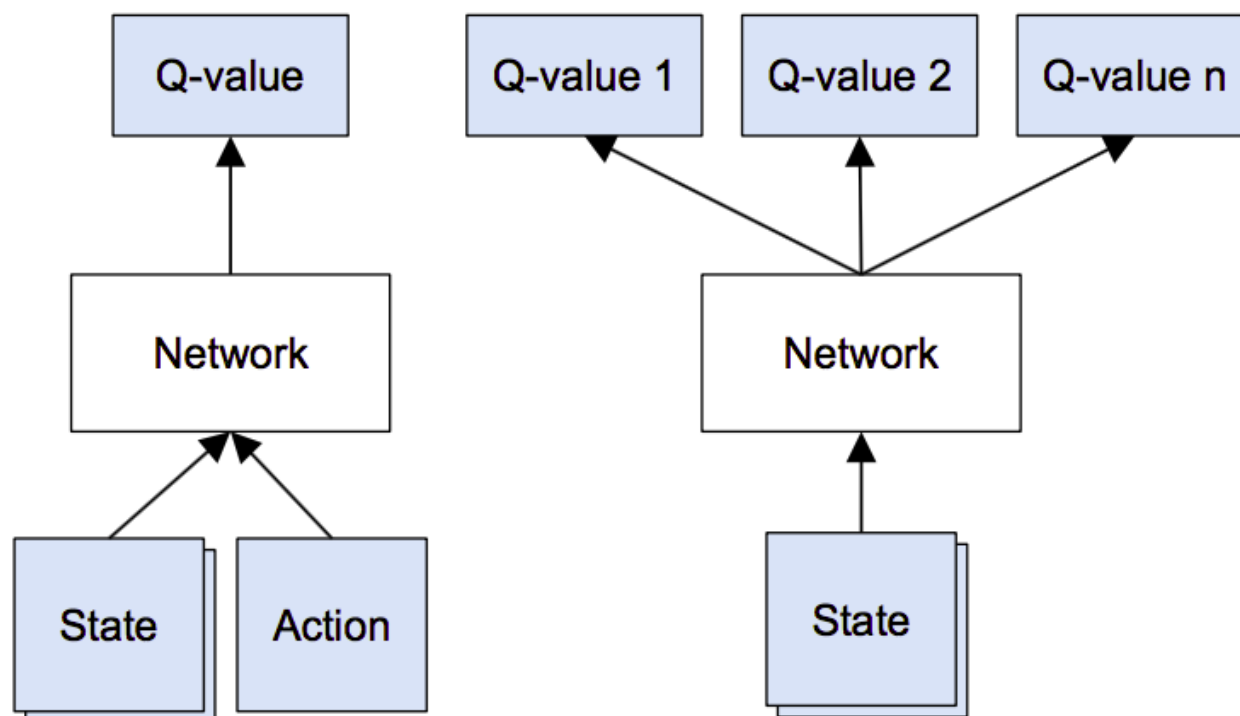
Deep learning approach: DQN



MDP again




Deep learning approach: DQN



$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Approximate Q-learning



Q-values:

$$\hat{Q}(s_t, a_t) = r + \gamma \cdot \operatorname{argmax}_{a'} \hat{Q}(s_{t+1}, a')$$


Objective:

$$L = (Q(s_t, a_t) - [r + \gamma \cdot \operatorname{argmax}_{a'} Q(s_{t+1}, a')])^2$$

Gradient step:

$$w_{t+1} = w_t - \alpha \cdot \frac{\delta L}{\delta w}$$

Approximate Q-learning



Q-values:


$$\hat{Q}(s_t, a_t) = r + \gamma \cdot \underset{a'}{\operatorname{argmax}} \hat{Q}(s_{t+1}, a')$$

Objective:

$$L = \left(Q(s_t, a_t) - \underbrace{\left[r + \gamma \cdot \underset{a'}{\operatorname{argmax}} Q(s_{t+1}, a') \right]}_{\text{consider const}} \right)^2$$

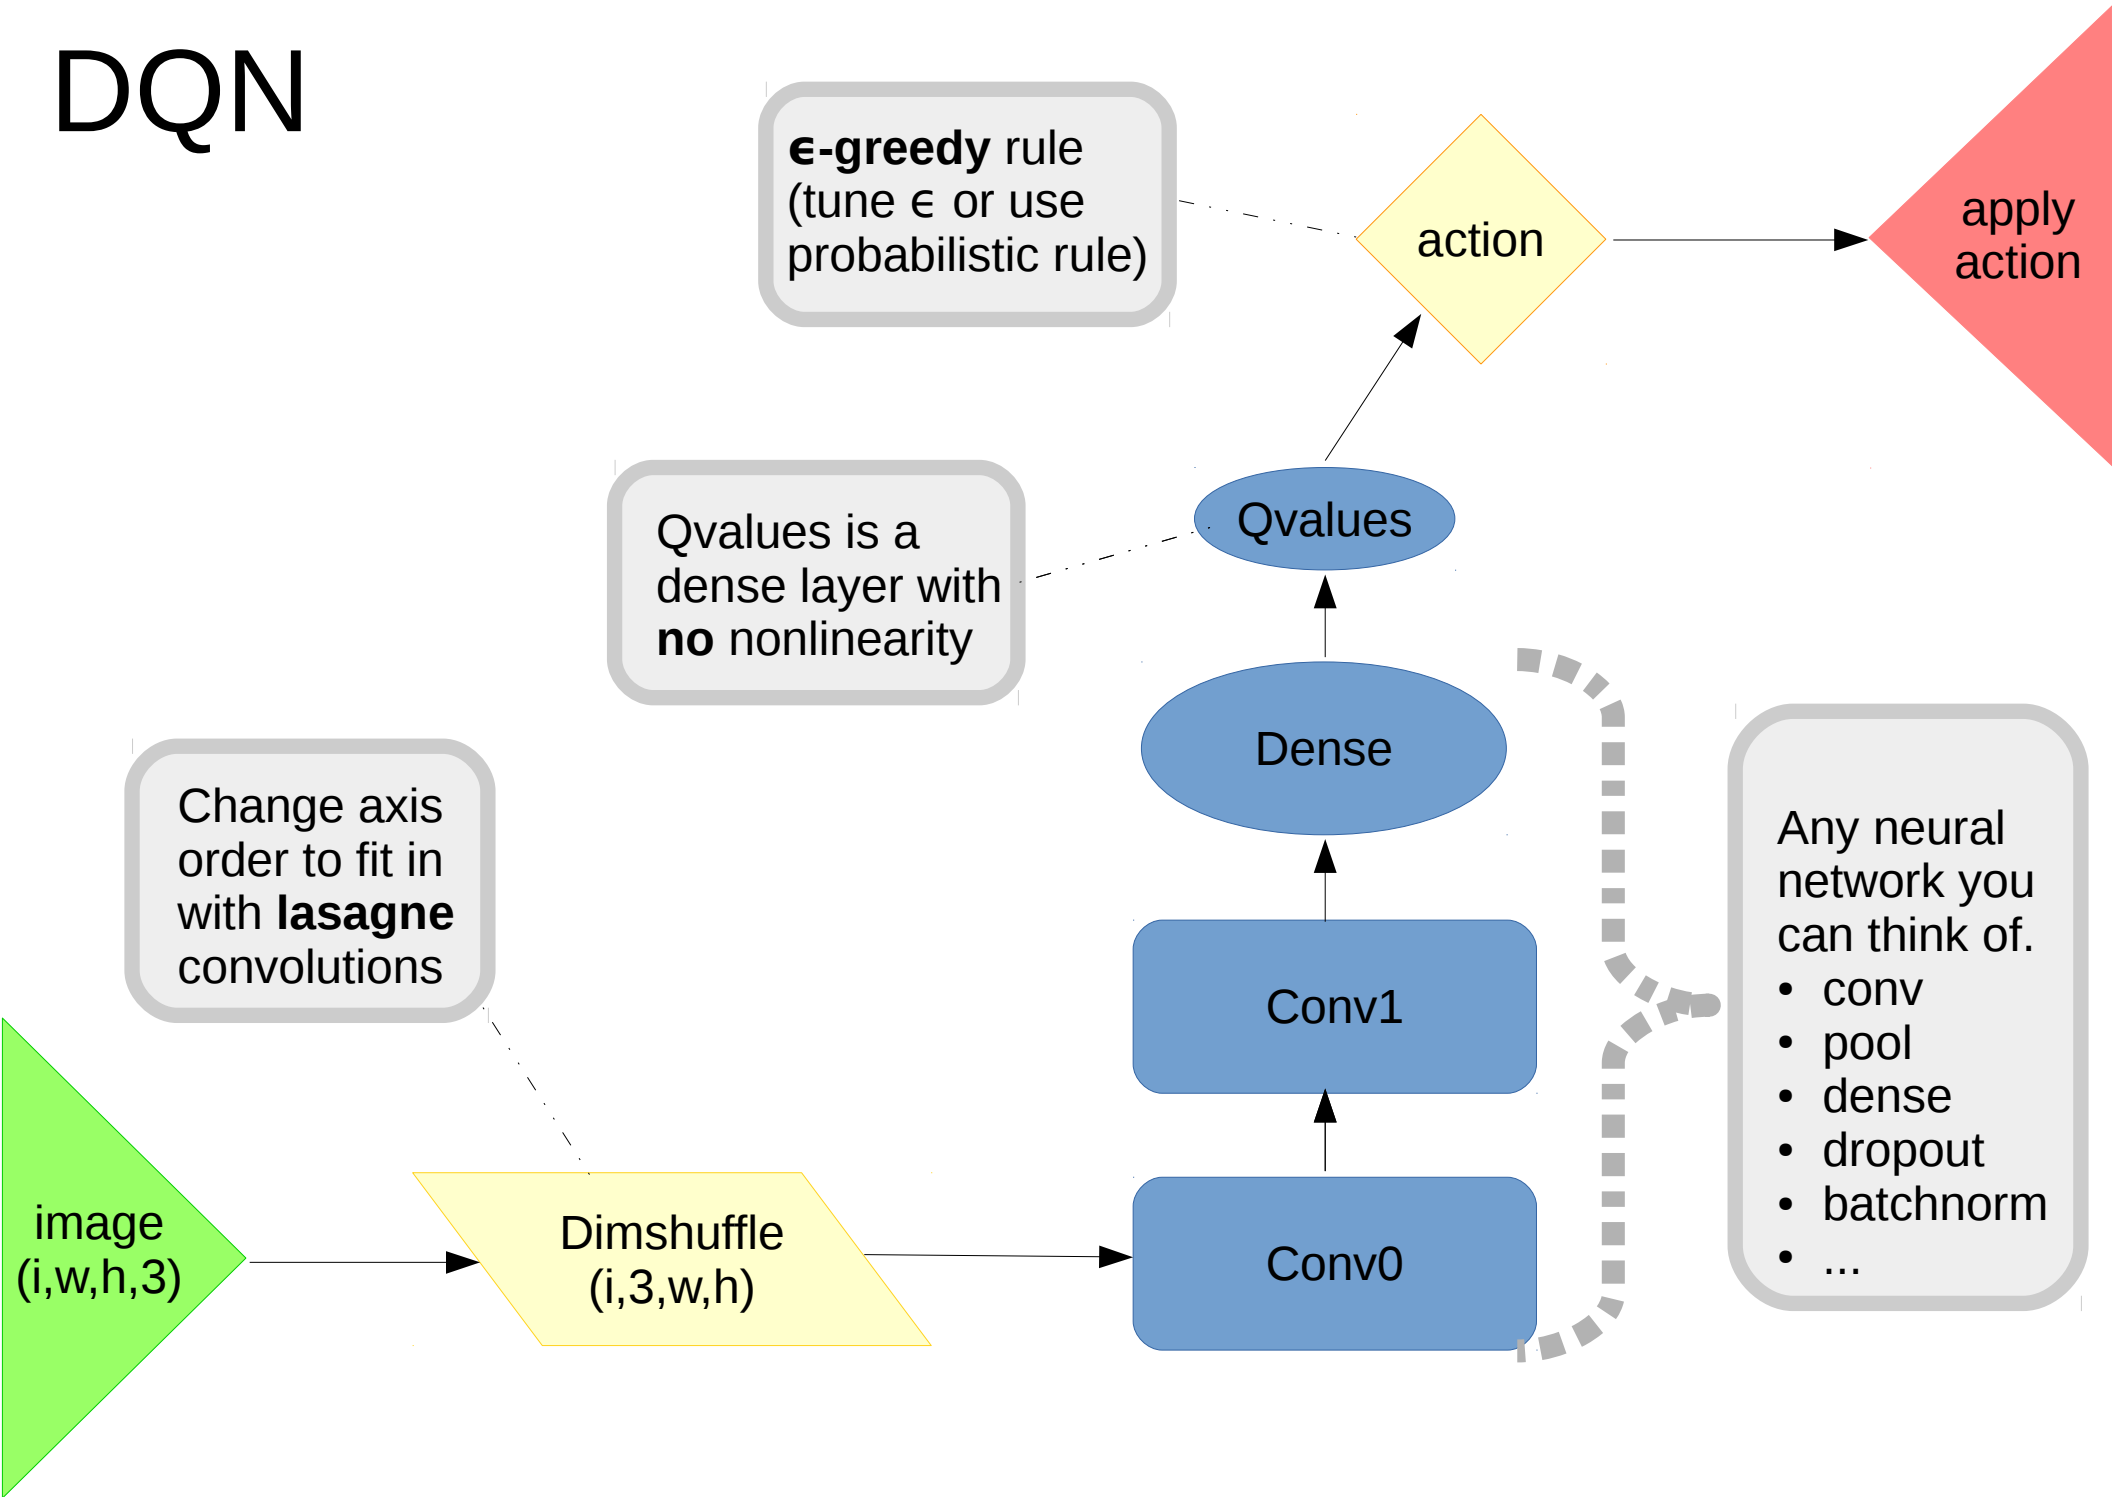
Gradient step:

$$w_{t+1} = w_t - \alpha \cdot \frac{\delta L}{\delta w}$$




Let's devise some network


DQN




Because TSNE



- Embedding of pre-last layer activations
- Color = state value = $\max_a Q(s,a)$





How bad is it if agent spent
last 1000 ticks under the left rock?


Problem

- Training samples are **not** “i.i.d”,
- Model forgets parts of environment it hasn't visited for some time
- Drops on learning curve
- **Any ideas?**



Multiple agent trick


Idea: Throw in several agents with shared \mathbf{W} .



Multiple agent trick

Idea: Throw in several agents with shared W .


- Chances are, they will be exploring different parts of the environment,
- More stable training,
- Requires a lot of interaction,
- Alternative to experience replay.



Experience replay

Idea: store several past interactions
 $\langle s, a, r, s' \rangle$
Train on random subsamples

Any +/- ?




Experience replay

Idea: store several past interactions
 $\langle s, a, r, s' \rangle$

Train on random subsamples

- Atari DQN: $>10^5$ interactions
- Closer to i.i.d
pool contains several sessions
- Older interactions were obtained under weaker policy
- Further development: prioritizing samples based on their importance.

<https://arxiv.org/abs/1511.05952>



Autocorrelation

- Reference is based on predictions

$$r + \gamma \cdot \operatorname{argmax}_{a'} Q(s_{t+1}, a')$$

- Any error in Q approximation is propagated to neighbors
- If some $Q(s,a)$ is mistakenly over-exaggerated, neighboring qvalues will also be increased in a cascade
- Worst case: divergence
- **Any ideas?**

Target networks

Idea: use older network snapshot
to compute reference

$$L = \left(Q(s_t, a_t) - \left[r + \gamma \cdot \underset{a'}{\operatorname{argmax}} Q^{\text{old}}(s_{t+1}, a') \right] \right)^2$$

- Update Q old periodically
 - Slows down training

Target networks

Idea: use older network snapshot
to compute reference

$$L = \left(Q(s_t, a_t) - \left[r + \gamma \cdot \operatorname{argmax}_{a'} Q^{\text{old}}(s_{t+1}, a') \right] \right)^2$$

- Update Q old periodically
 - Slows down training
- Smooth version:
 - use moving average

$$\theta^{\text{old}} := (1 - \alpha) \cdot \theta^{\text{old}} + \alpha \cdot \theta^{\text{new}}$$

- Θ = weights

Final problem



Left or right?

Problem:

Most practical cases are partially observable:

Agent observation does not hold all information about process state
(e.g. human field of view).

Any ideas?

Problem:

Most practical cases are partially observable:

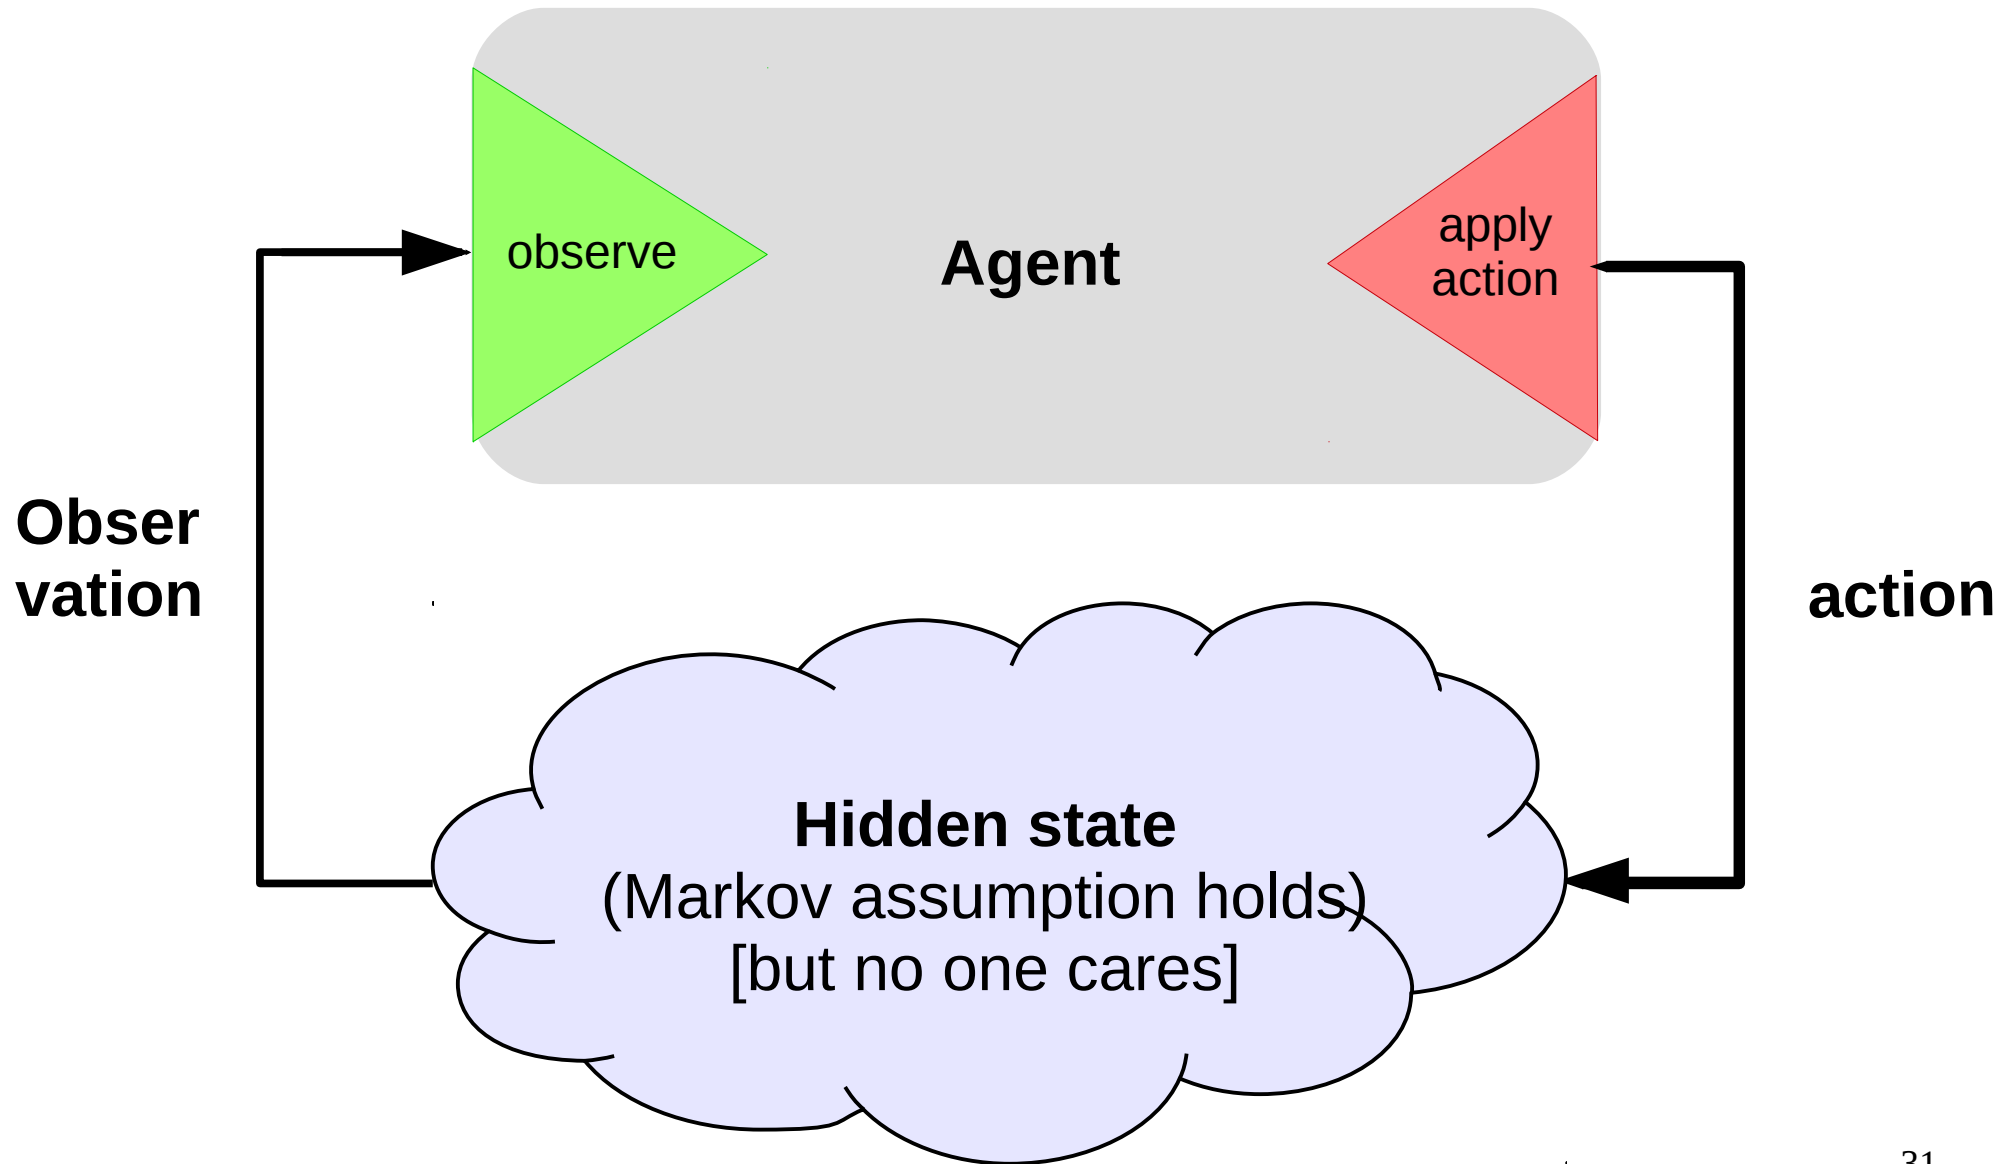
Agent observation does not hold all information about process state
(e.g. human field of view).

- However, we can try to infer hidden states from sequences of observations.

$$s_t \simeq m_t : P(m_t | o_t, m_{t-1})$$

- Intuitively that's agent memory state.

Partially observable MDP




N-gram heuristic

Idea:

$$s_t \neq o(s_t)$$

$$s_t \approx (o(s_{t-n}), a_{t-n}, \dots, o(s_{t-1}), a_{t-1}, o(s_t))$$

e.g. ball movement in breakout

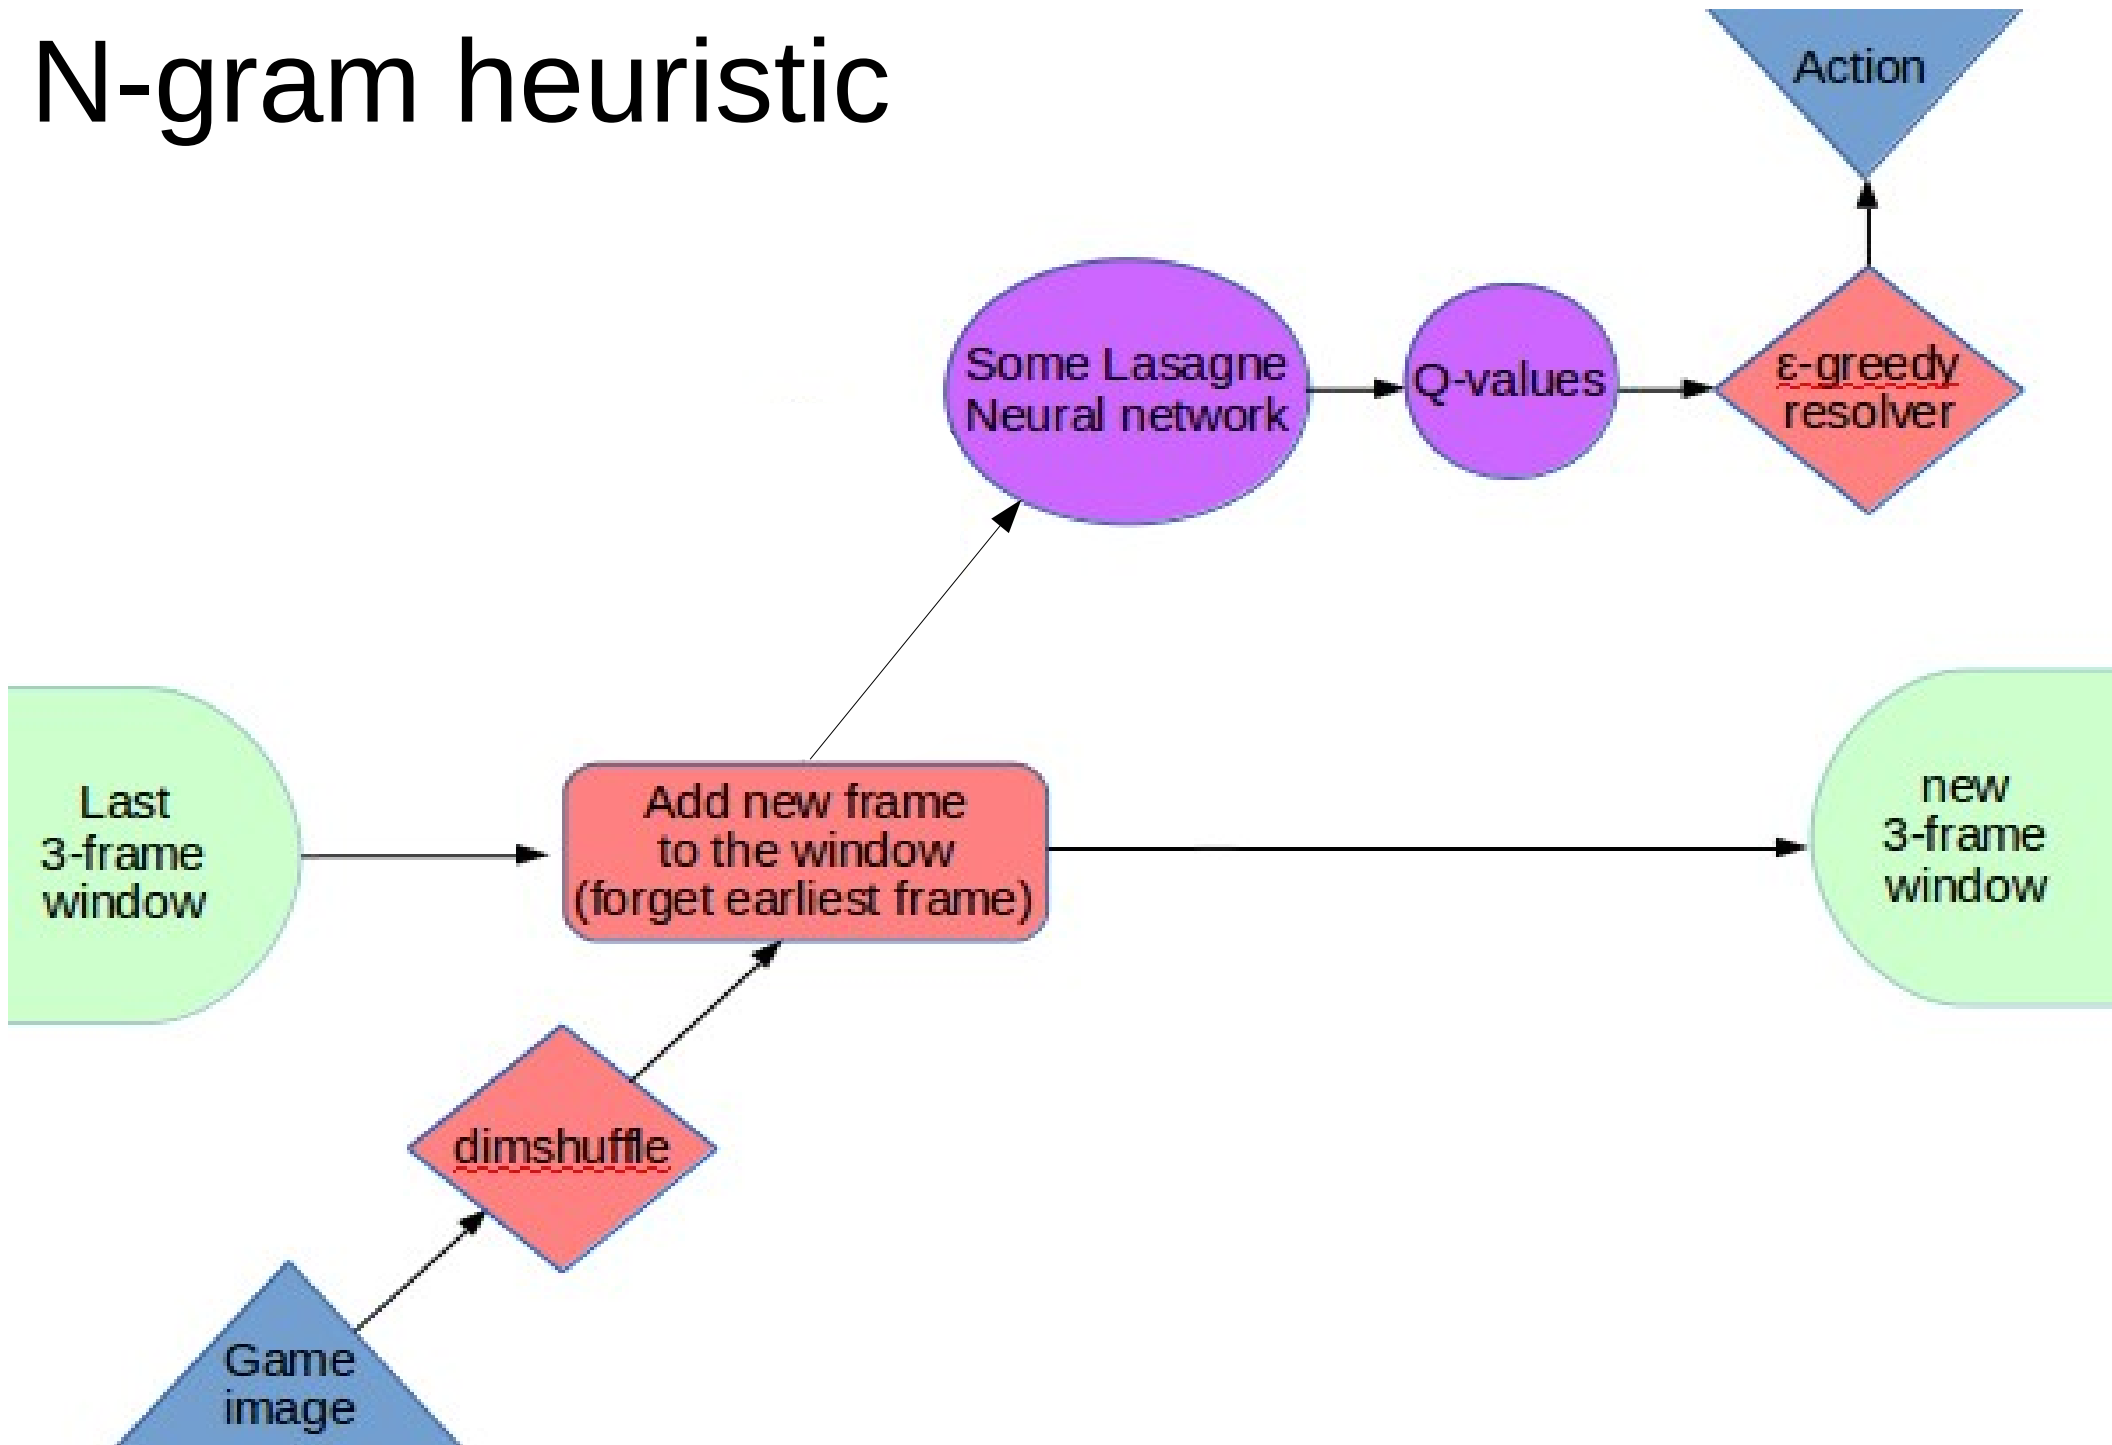


• One frame

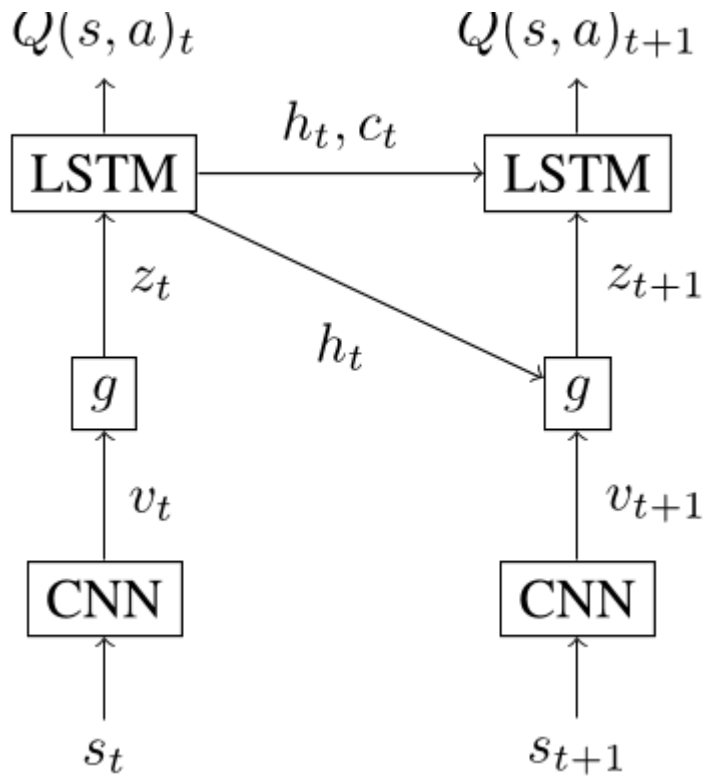


• Several frames

N-gram heuristic






Deep Recurrent RL




Recurrent agent memory

- Agent has his own hidden state.
- Trained via BPTT with a fixed depth
- Problem: next input depends on chosen action
- Even more autocorrelations :)




Deep Recurrent RL

Learning curves for KungFuMaster



CHALLENGE ACCEPTED



LET'S CODE IT

memegenerator.net

http://mybinder.org/repo/yandexdataschool/dqn_binder

Most important slide

RL isn't magical

- It won't learn everything in the world given any data
- Requires interaction
- Sparse and/or delayed rewards are a major problem
- Less playing Atari, more real world problems
No, doom is not a real world problem, dummy!
- Getting rid of heuristics towards mathematical soundness
- Machine Intelligence revolution date TBA